Lecture 4: Linear Regression and Classification

STATS 202: Statistical Learning and Data Science

Linh Tran

tranlm@stanford.edu

Department of Statistics
Stanford University

July 2, 2025

Announcements

- No section this Friday
- HW1 due tomorrow
  - Can ask for regrades up to a week from grades being released
  - Solutions will be posted next week
- Accommodation requests for midterms (in 2 weeks)

Outline

- Regression issues
- Comparing linear regression to KNN
- More classification
  - Logistic regression
  - Linear/quadratic discriminant analysis

Recap

Potential issues in linear regression

- Interactions between predictors
- Non-linear relationships
- Correlation of error terms
- Non-constant variance of error (heteroskedasticity)
- Outliers
- High leverage points
- Collinearity
- Mis-specification

Potential issues in linear regression

- Interactions between predictors
- Non-linear relationships
- Correlation of error terms
- Non-constant variance of error (heteroskedasticity)
- Outliers
- High leverage points
- Collinearity
- Mis-specification

Interactions between predictors

Linear regression has an *additive* assumption, e.g.:

$$\textit{sales} = \beta_0 + \beta_1 \cdot \texttt{tv} + \beta_2 \cdot \texttt{radio} + \epsilon \tag{1}$$

e.g. An increase of \$ 100 dollars in TV ads correlates to a fixed increase in sales, independent of how much you spend on radio ads.

If we visualize the residuals, it is clear that this is false:

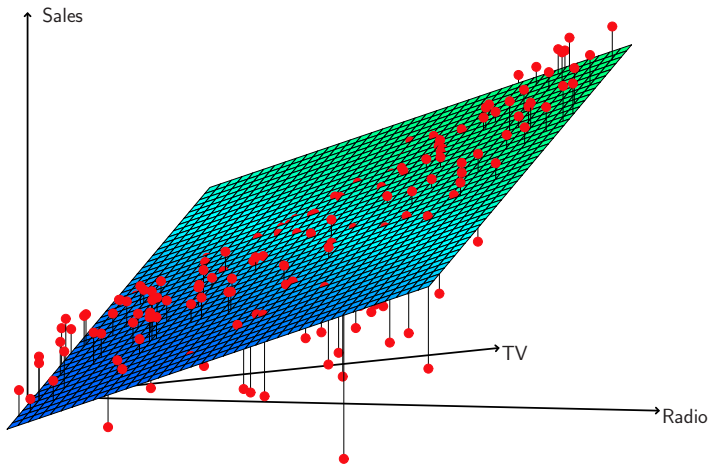

Interactions between predictors

One way to deal with this:

- Include multiplicative variables (aka interaction variables) in the model

$$sales = \beta_0 + \beta_1 \cdot tv + \beta_2 \cdot radio + \beta_3 \cdot (tv \times radio) + \epsilon \quad (2)$$

- Makes the effect of TV ads dependent on the radio ads (and vice versa)
- The *interaction variable* is high when both tv and radio are high

Interactions between predictors

Two ways of including interaction variables (in R):

- Create a new variable that is the product of the two
- Specify the interaction in the model formula

```
> lm.fit=lm(Sales~.+Income:Advertising+Price:Age,data=Carseats)
> summary(lm.fit)

Call:
lm(formula = Sales ~ . + Income:Advertising + Price:Age, data =
    Carseats)

Residuals:
    Min      1Q  Median      3Q     Max
-2.921  -0.750   0.018   0.675   3.341

Coefficients:
                    Estimate  Std. Error t value Pr(>|t|)
(Intercept)         6.575565    1.008747    6.52  2.2e-10 ***
CompPrice           0.092937    0.004118   22.57  < 2e-16 ***
Income              0.010894    0.002604    4.18  3.6e-05 ***
Advertising         0.070246    0.022609    3.11  0.00203 **
Population          0.000159    0.000368    0.43  0.66533
Price              -0.100806    0.007440  -13.55  < 2e-16 ***
ShelveLocGood       4.848676    0.152838   31.72  < 2e-16 ***
ShelveLocMedium     1.953262    0.125768   15.53  < 2e-16 ***
Age                -0.057947    0.015951   -3.63  0.00032 ***
Education          -0.020852    0.019613   -1.06  0.28836
UrbanYes            0.140160    0.112402    1.25  0.21317
USYes              -0.157557    0.148923   -1.06  0.29073
Income:Advertising  0.000751    0.000278    2.70  0.00729 **
Price:Age           0.000107    0.000133    0.80  0.42381
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Non-linear relationships

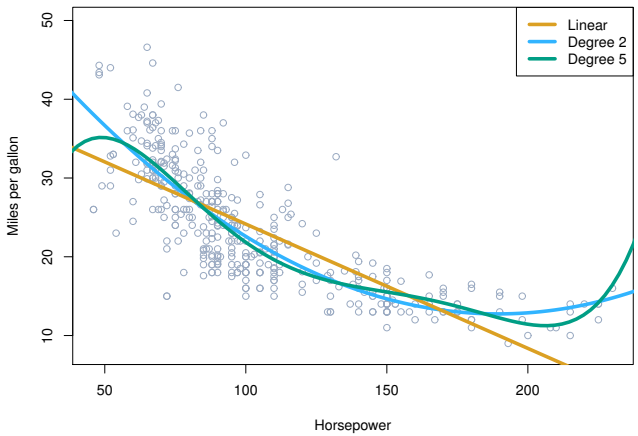

Scatterplots between X and Y may reveal non-linear relationships

- **Solution**: Include polynomial terms in the model

$$\begin{aligned} MPG = &\beta_0 + \beta_1 \cdot horsepower \\ &+ \beta_2 \cdot horsepower^2 \\ &+ \beta_3 \cdot horsepower^3 + ... + \epsilon \end{aligned} \tag{3}$$

Non-linear relationships

In 2 or 3 dimensions, this is easy to visualize. What do we do when we have too many predictors?

Non-linear relationships

In 2 or 3 dimensions, this is easy to visualize. What do we do when we have too many predictors?

Plot the residuals against the response and look for a pattern:

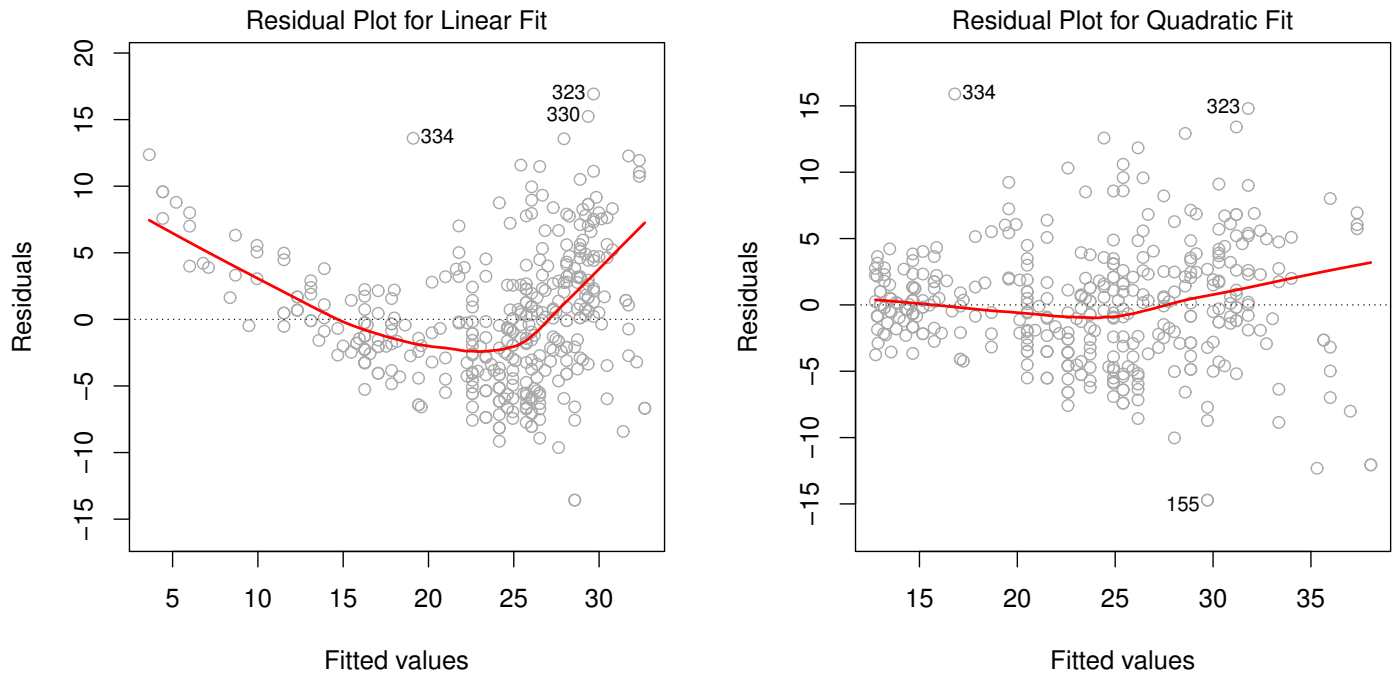

Correlation of error terms

We assumed that the errors for each sample are independent:

$$y_i = f(x_i) + \epsilon_i : \epsilon_i \overset{iid}{\sim} \mathcal{N}(0, \sigma^2) \tag{4}$$

Correlation of error terms

We assumed that the errors for each sample are independent:

$$y_i = f(x_i) + \epsilon_i : \epsilon_i \overset{iid}{\sim} \mathcal{N}(0, \sigma^2) \tag{4}$$

When it doesn't hold:

- Invalidates any assertions about Standard Errors, confidence intervals, and hypothesis tests

Example: Suppose that by accident, we double the data (i.e. we use each sample twice). Then, the standard errors would be artificially smaller by a factor of $\sqrt{2}$.

Correlation of error terms

Examples of when this happens:

- *Time series*: Each sample corresponds to a different point in time. The errors for samples that are close in time are correlated.
- *Spatial data*: Each sample corresponds to a different location in space.
- *Clustered data*: Study on predicting height from weight at birth. Suppose some of the subjects in the study are in the same family, their shared environment could make them deviate from $f(x)$ in similar ways.

Correlation of error terms

Simulations of time series with increasing correlations on ϵ_i.

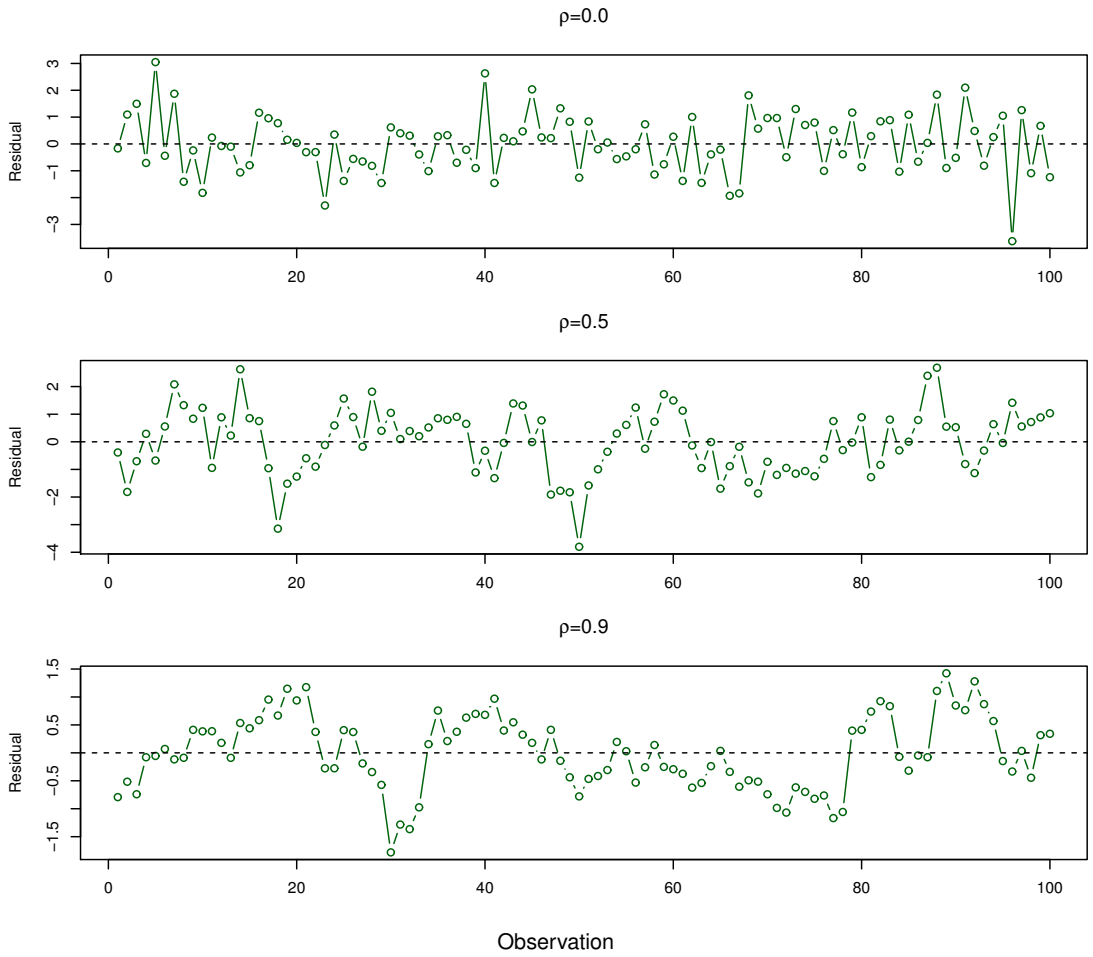

Non-constant variance of error (heteroskedasticity)

The variance of the error depends on the input value.

To diagnose this, we can plot residuals vs. fitted values:

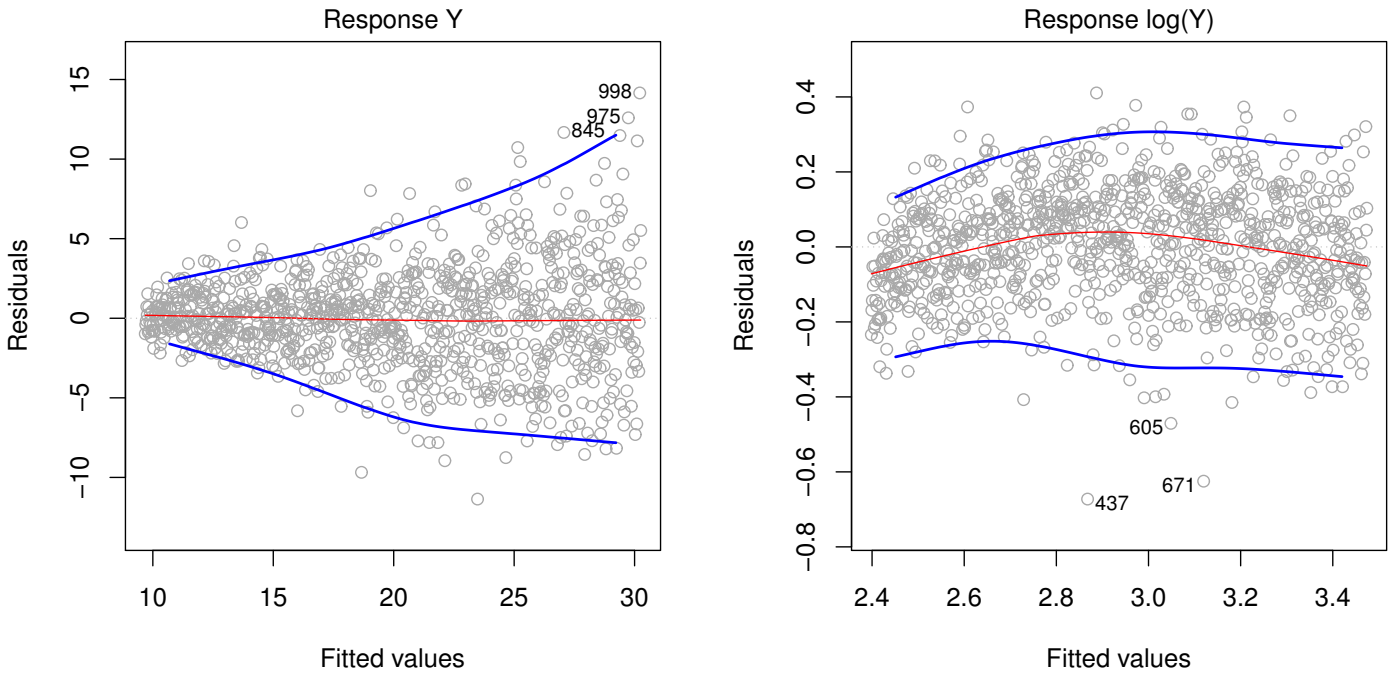

Non-constant variance of error (heteroskedasticity)

The variance of the error depends on the input value.

To diagnose this, we can plot residuals vs. fitted values:

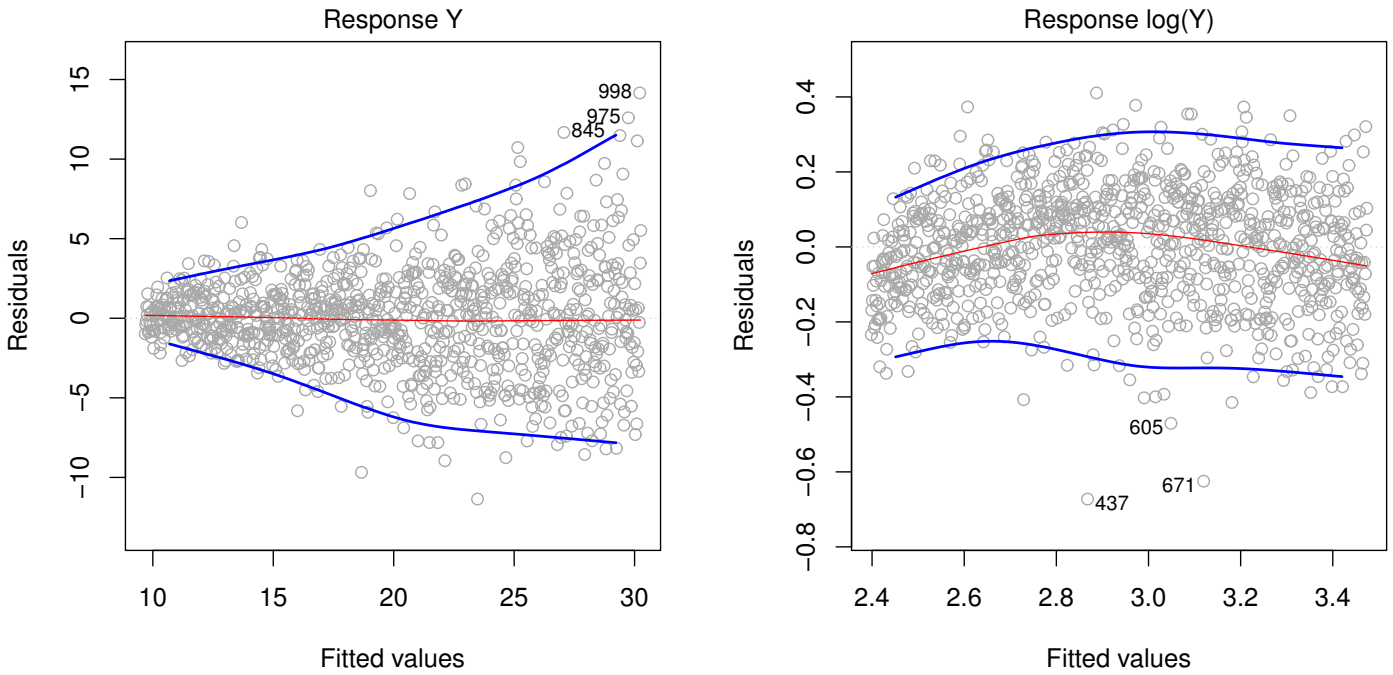

Solution: If the trend in variance is relatively simple, we can transform the response using a logarithm, for example.

Outliers

Outliers are points with very large errors, e.g.

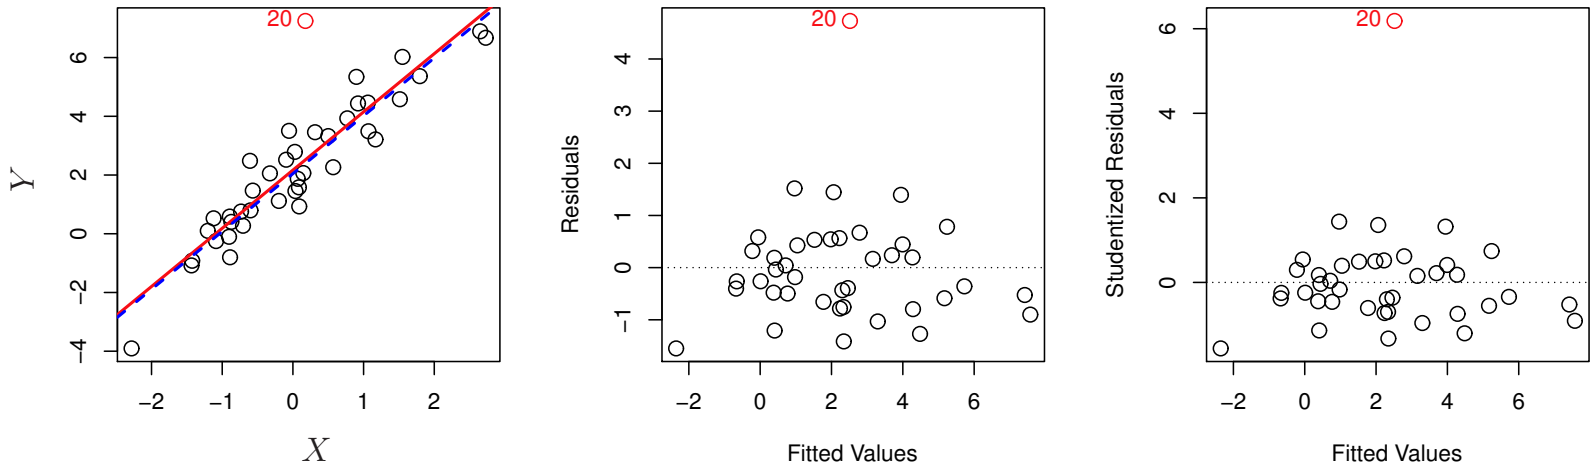

While they may not affect the fit, they might affect our assessment of model quality.

Outliers

Outliers are points with very large errors, e.g.

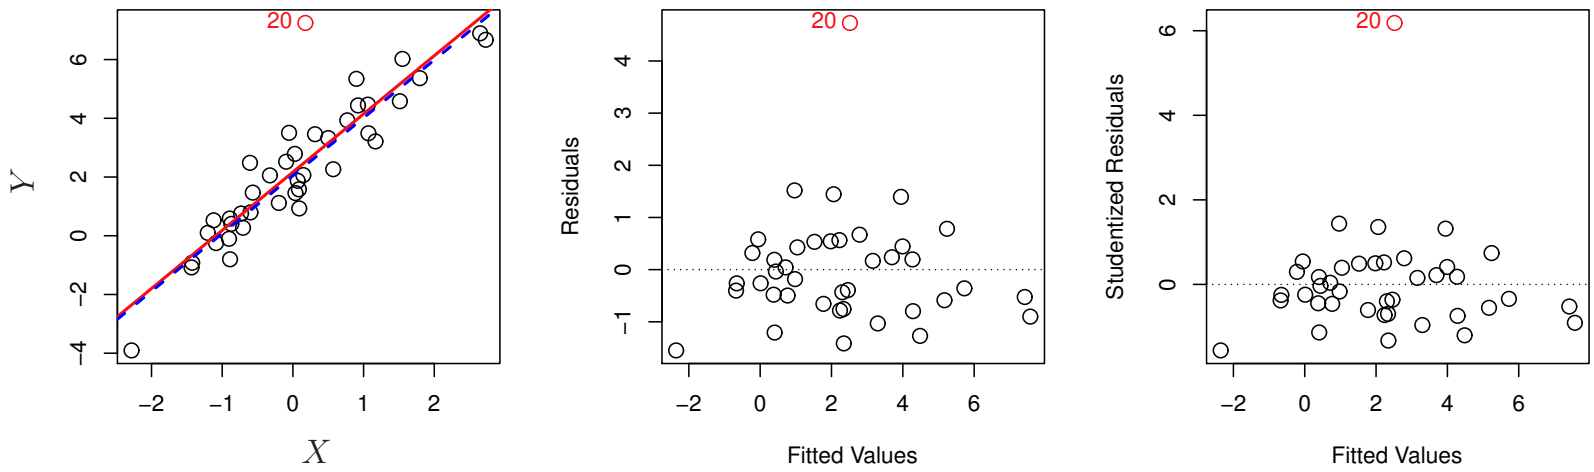

While they may not affect the fit, they might affect our assessment of model quality.

Possible solutions:

- If we believe an outlier is due to an error in data collection, we can remove it.

High leverage points

Some samples with extreme inputs have a large effect on $\hat{\beta}$.

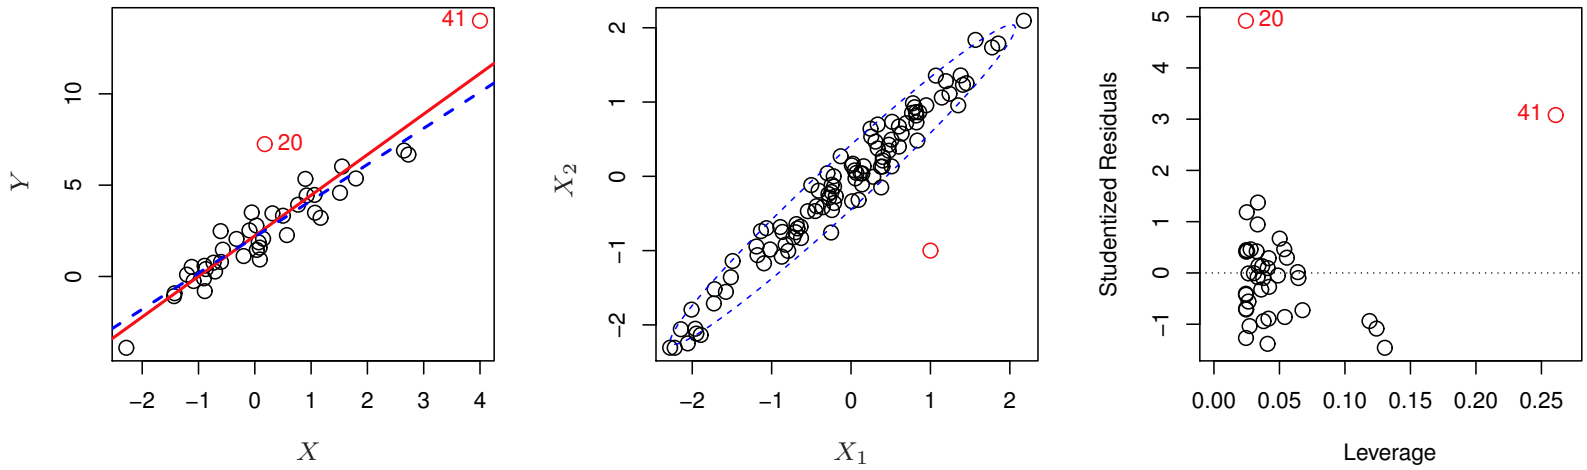

This can be measured with the *leverage statistic* or *self influence*:

$$h_{ii} = \frac{\partial \hat{y}_i}{\partial y_i} = (\underbrace{\mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1}\mathbf{X}^\top}_{\text{Hat matrix}})_{i,i} \in \left[\frac{1}{n}, 1\right] \tag{5}$$

Values closer to 1 have high leverage.

Studentized residuals

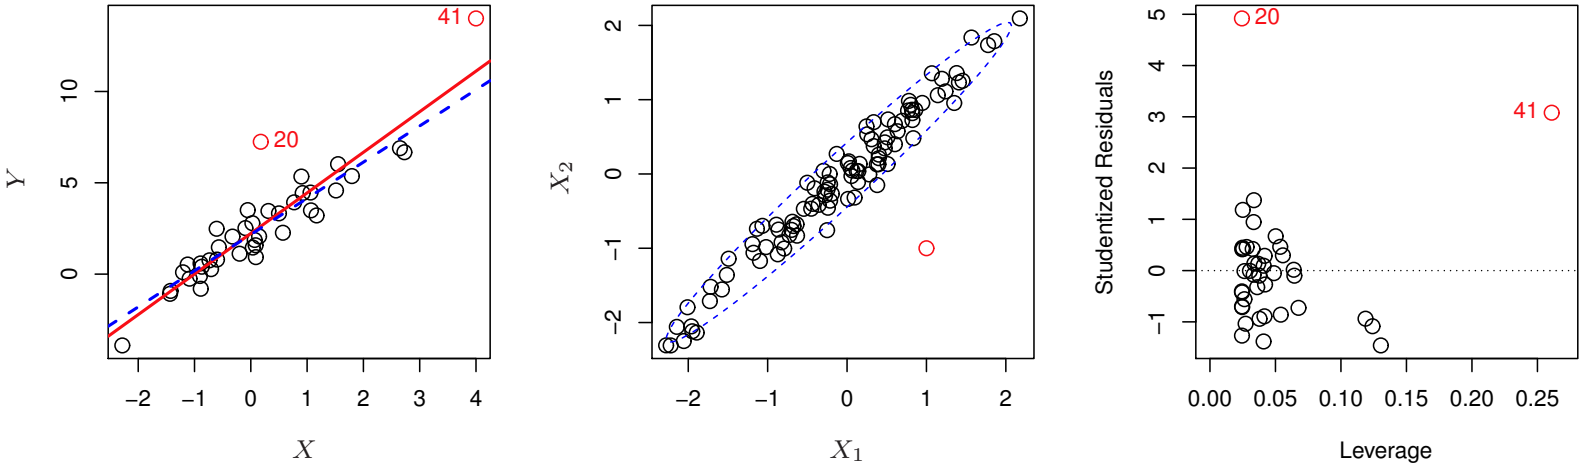

- The residual $\hat{\epsilon}_i = y_i - \hat{y}_i$ is an estimate for the noise ϵ_i
- The standard error of $\hat{\epsilon}_i$ is $\sigma\sqrt{1 - h_{ii}}$
- A studentized residual is $\hat{\epsilon}_i$ divided by its standard error
- It follows a Student-t distribution with $n - p - 2$ degrees of freedom

Collinearity

Two predictors are collinear if one explains the other well, e.g.

$$limit = a \times rating + b \tag{6}$$

i.e. they contain the same information

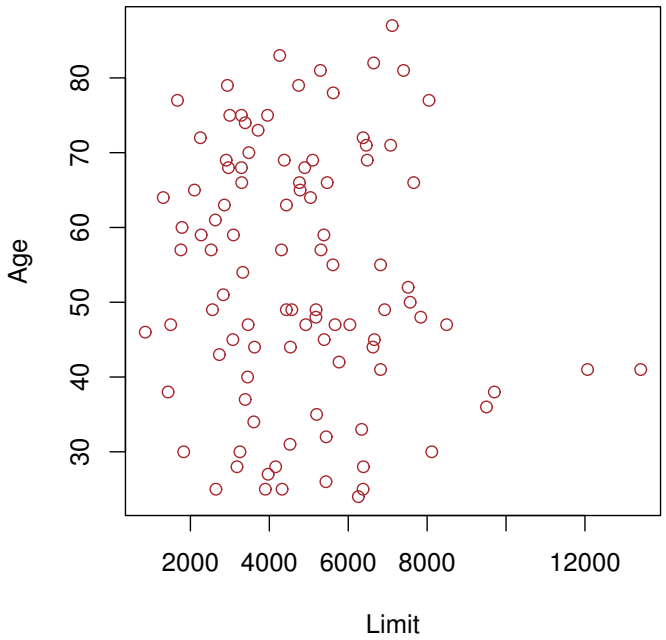

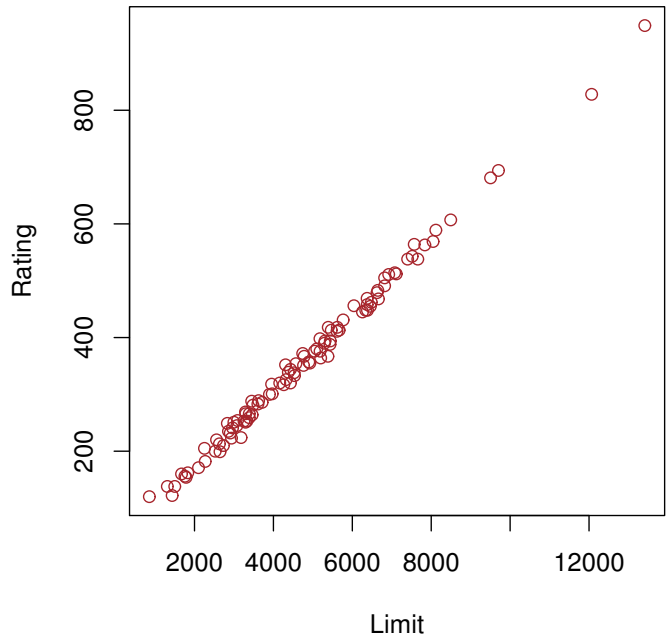

Collinearity

Problem: The coefficients become *unidentifiable*.

- i.e. different coefficients can mean the same fit

Collinearity

Problem: The coefficients become *unidentifiable*.

- i.e. different coefficients can mean the same fit

Example: using two identical predictors (`limit`):

$$\begin{aligned} \textit{balance} &= \beta_0 + \beta_1 \cdot \textit{limit} + \beta_2 \cdot \textit{limit} && (7) \\ &= \beta_0 + (\beta_1 + 100) \cdot \textit{limit} + (\beta_2 - 100) \cdot \textit{limit} && (8) \end{aligned}$$

The fit $(\hat{\beta}_0, \hat{\beta}_1, \hat{\beta}_2)$ is just as good as $(\hat{\beta}_0, \hat{\beta}_1 + 100, \hat{\beta}_2 - 100)$

Collinearity effect

Collinearity results in unstable estimates of β.

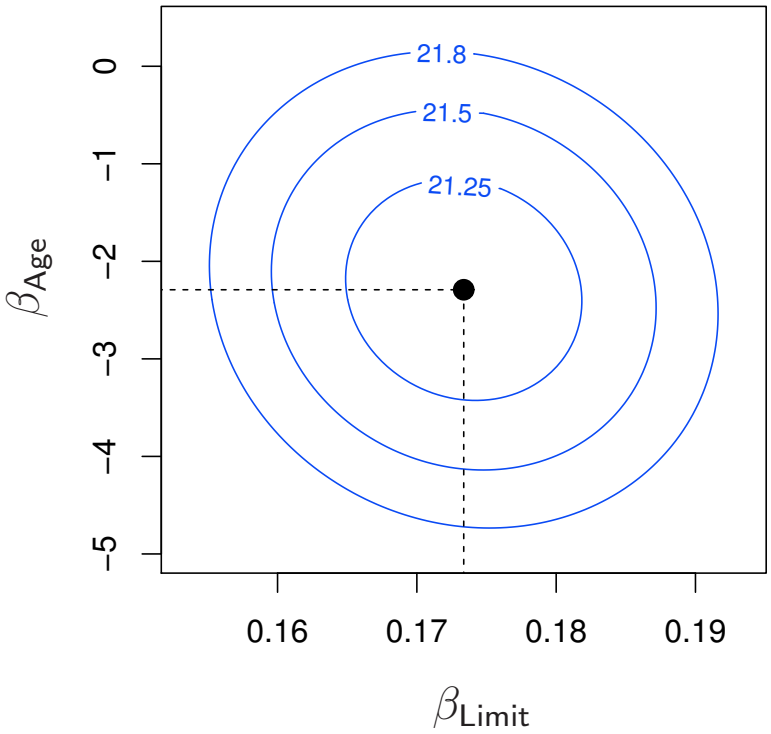

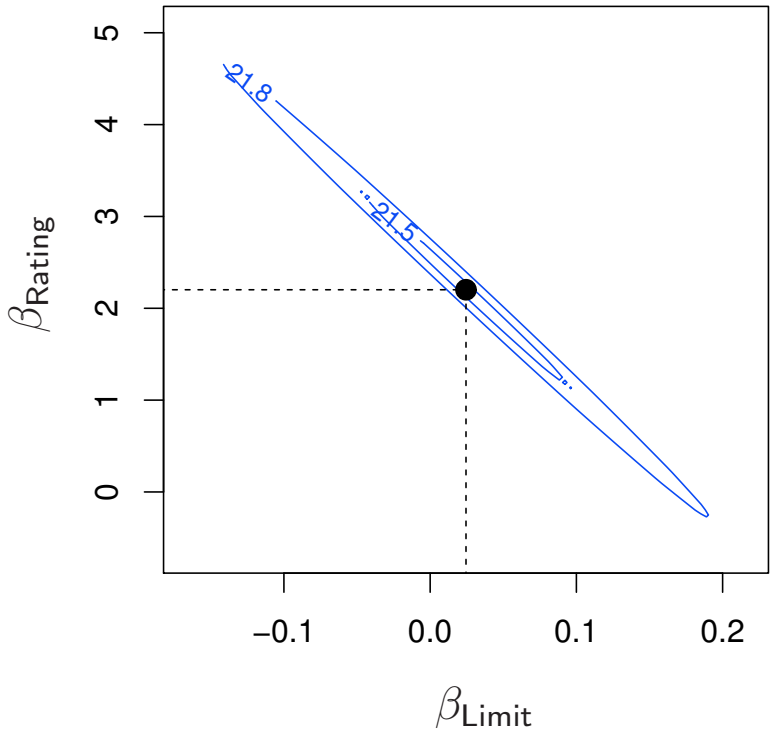

Recognizing collinearity

If 2 variables are collinear, we can easily diagnose this using their correlation.

A group of q variables is multilinear if these variables "contain less information" than q independent variables. Pairwise correlations may not reveal multilinear variables.

The Variance Inflation Factor (VIF) measures how necessary a variable is, or how predictable it is given the other variables:

$$VIF(\hat{\beta}_j) = \frac{1}{1 - R^2_{X_j|X_{-j}}}, \tag{9}$$

where $R^2_{X_j|X_{-j}}$ is the R^2 statistic for multiple linear regression of the predictor X_j onto the remaining predictors.

Dealing with collinearity

Three primary ways:

1. Drop one of the correlated features (e.g. Ridge/LASSO).
2. Combine the correlated features (e.g. PCA).
3. More data.

Mis-specification

What if our true distribution P_0 **isn't** linear?

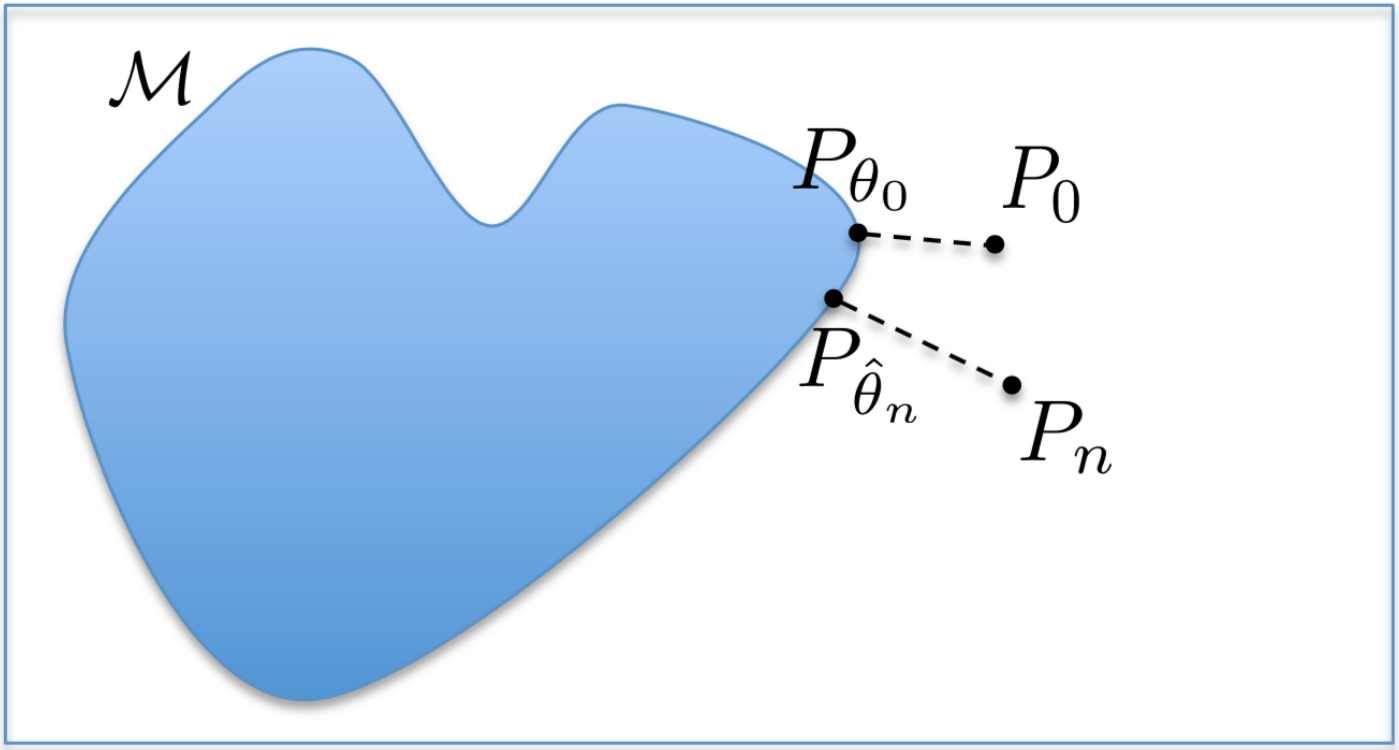

Estimates will still converge to a fixed value within our model, e.g.

$$\theta_0 \triangleq \arg\min_{\theta} D(P_\theta, P_0) : \theta = (\beta_0, ..., \beta_p) \tag{10}$$

Comparing Linear Regression to K-nearest neighbors

Linear regression: prototypical parametric method
KNN regression: prototypical nonparametric method

$$\hat{f}_n(x) = \frac{1}{K} \sum_{i \in N_K(x)} y_i \tag{11}$$

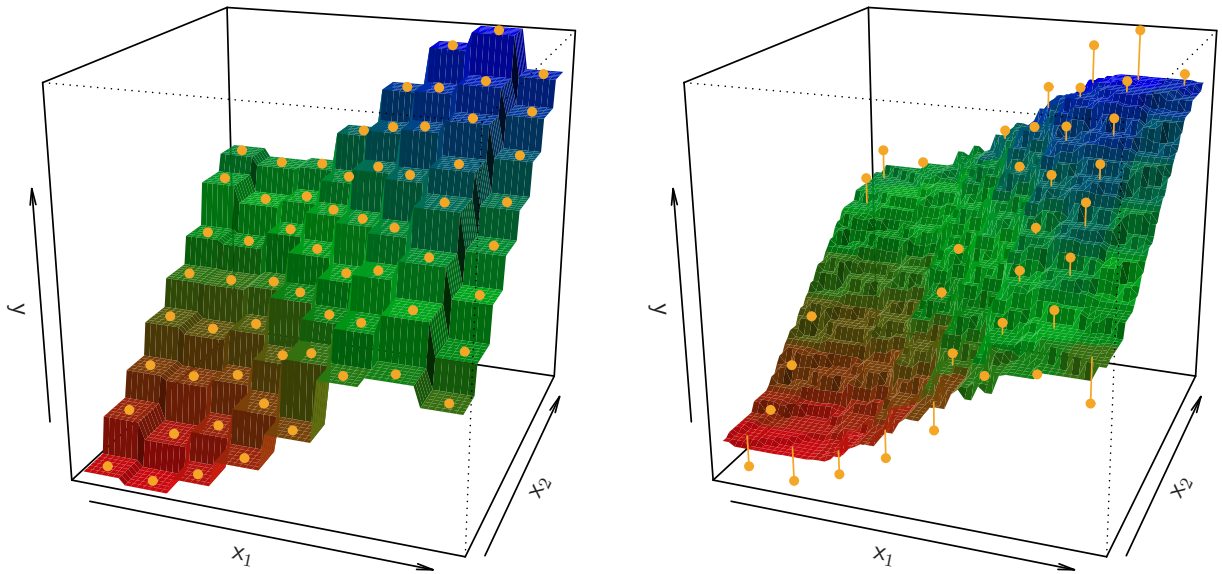

Examples of KNN with $K = 1$ (left) and $K = 9$ (right)

Comparing Linear Regression to K-nearest neighbors

Linear regression: prototypical parametric method
KNN regression: prototypical nonparametric method Long story short:

- KNN is better when the function f_0 is not linear (and plenty of data)
  - **Question**: What if the true function f_0 IS linear?

Comparing Linear Regression to K-nearest neighbors

Linear regression: prototypical parametric method
KNN regression: prototypical nonparametric method Long story short:

- KNN is better when the function f_0 is not linear (and plenty of data)
  - **Question**: What if the true function f_0 IS linear?
- When n is not much larger than p, even if f_0 is nonlinear, linear regression can outperform KNN.

Comparing Linear Regression to K-nearest neighbors

Linear regression: prototypical parametric method
KNN regression: prototypical nonparametric method Long story short:

- KNN is better when the function f_0 is not linear (and plenty of data)
  - **Question**: What if the true function f_0 IS linear?
- When n is not much larger than p, even if f_0 is nonlinear, linear regression can outperform KNN.
- KNN has smaller bias, but this comes at a price of (much) higher variance (c.f. overfitting)

Comparing Linear Regression to K-nearest neighbors

KNN estimates for a simulation from a linear model

- True function f_0 is linear

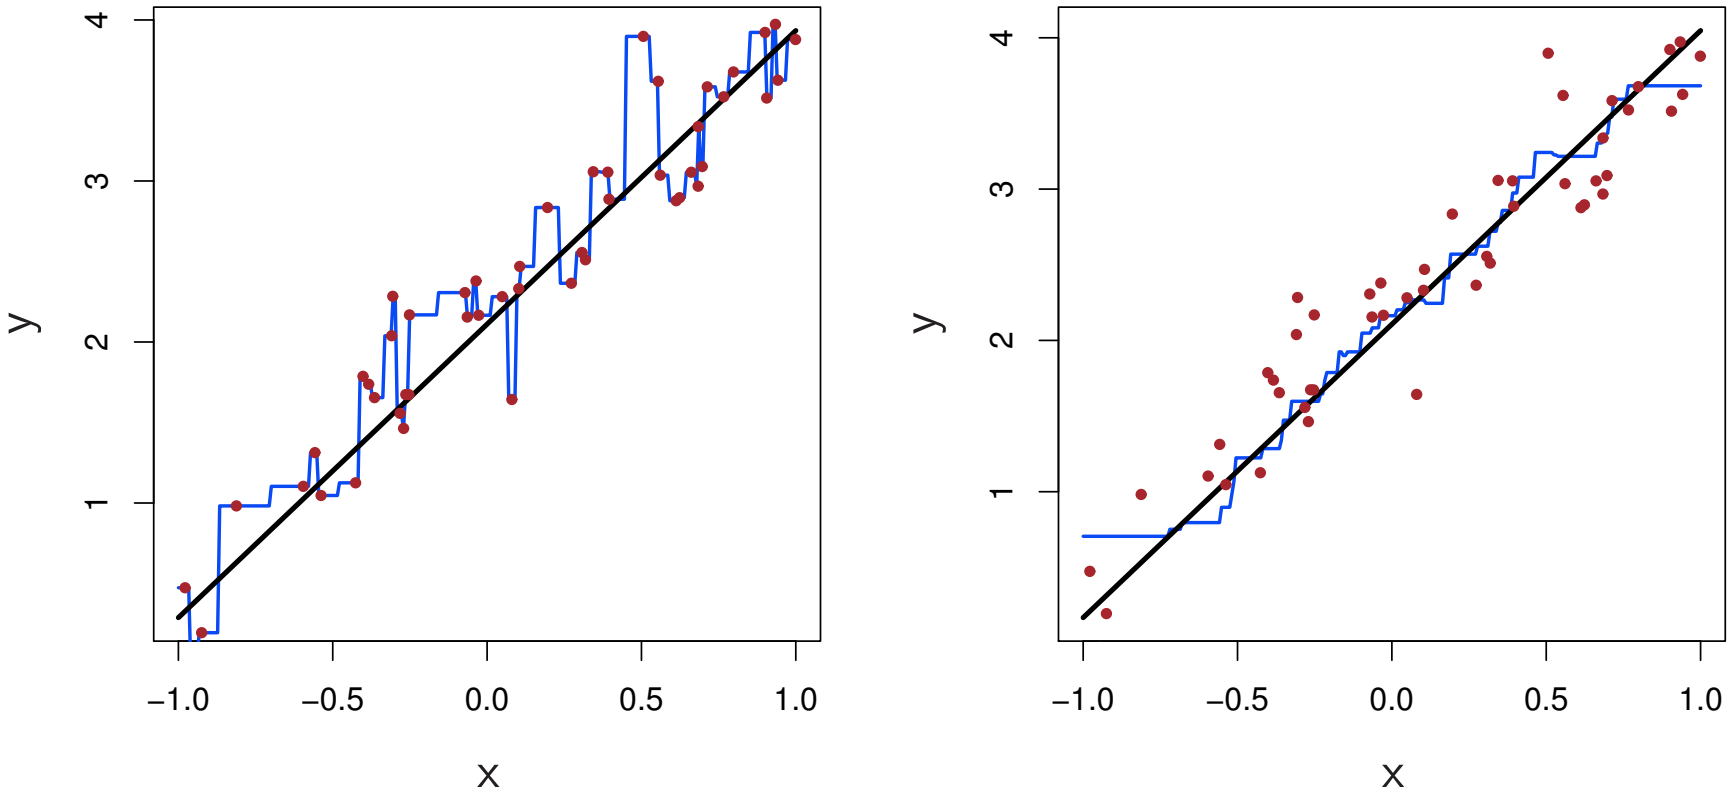

KNN fits with $K = 1$ (left) and $K = 9$ (right)

Comparing Linear Regression to K-nearest neighbors

Linear models dominate KNN

- We're able to gain statistical efficiency by taking advantage of the linear association

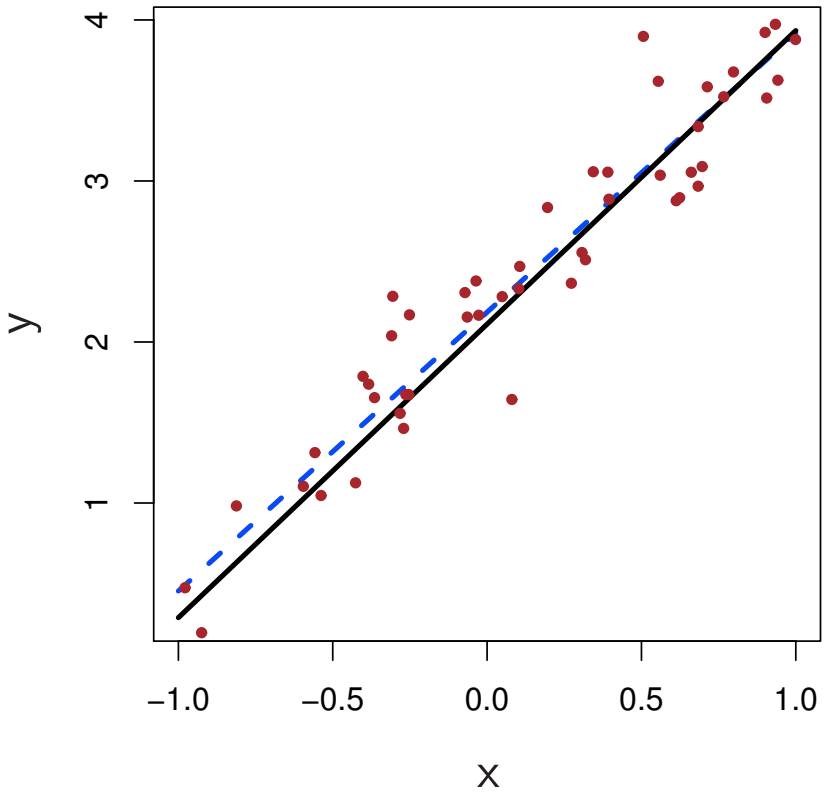

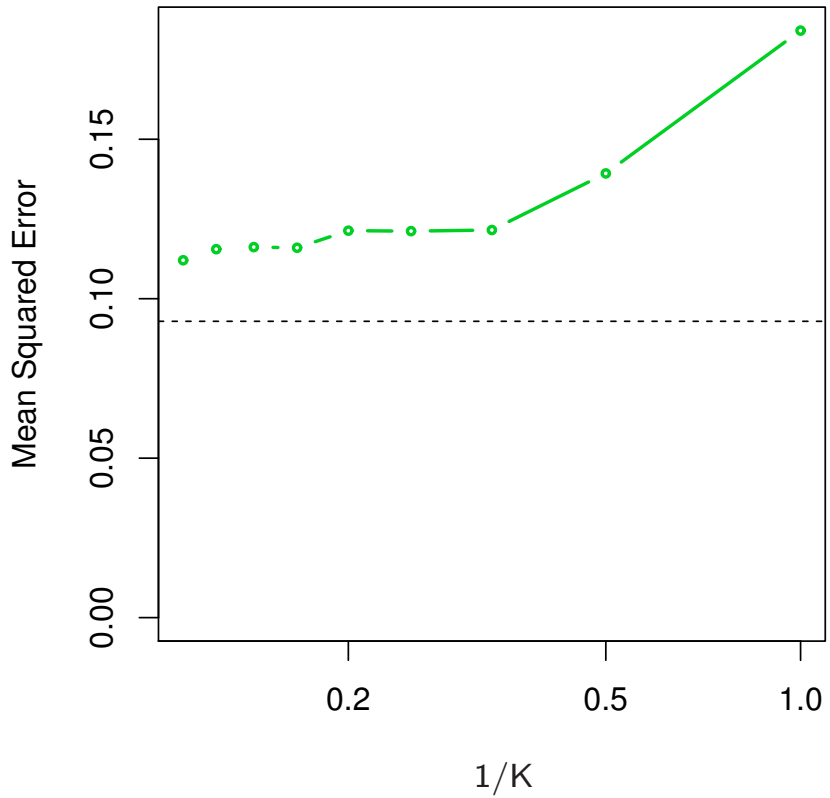

Comparing Linear Regression to K-nearest neighbors

Increasing deviations from linearity

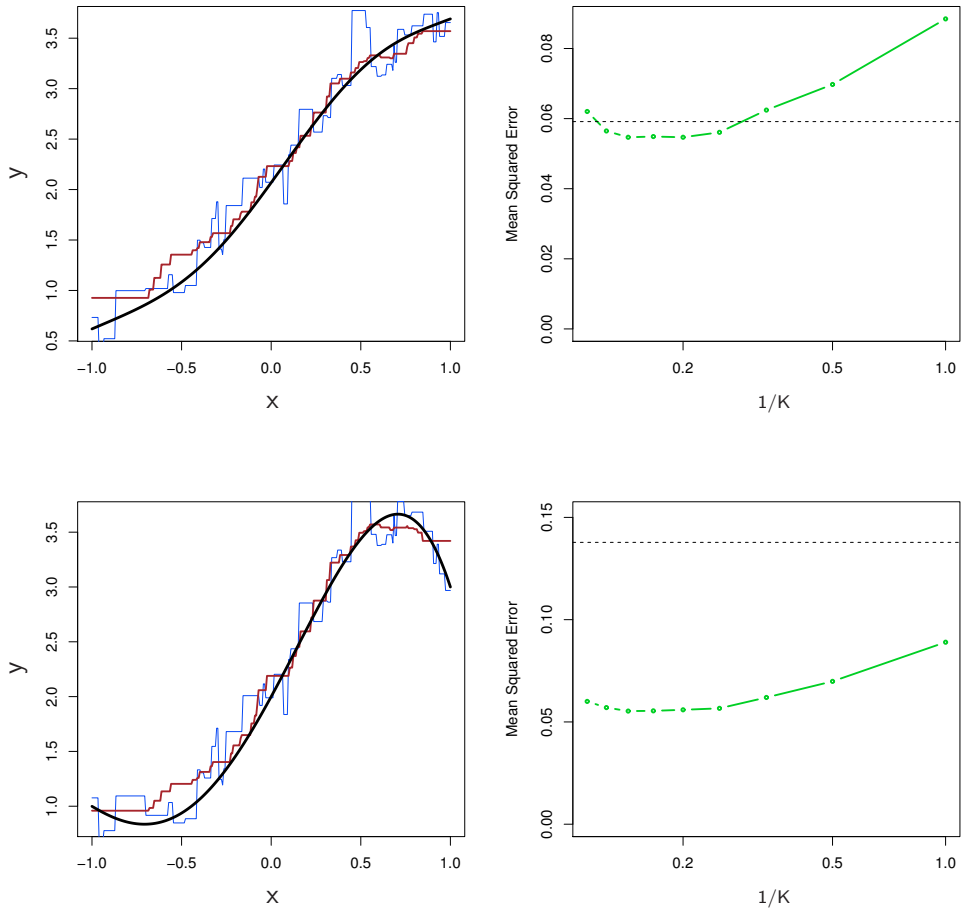

Comparing Linear Regression to K-nearest neighbors

When there are more predictors than observations, linear regression dominates

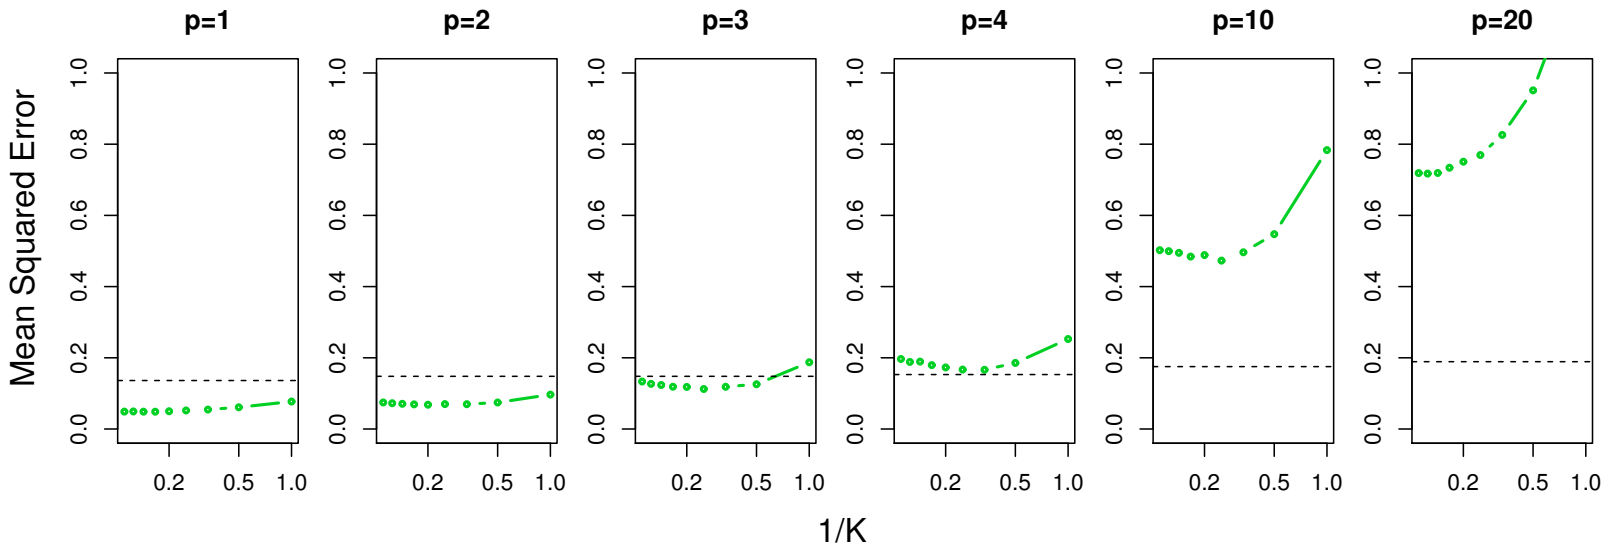

When $p >> n$, each sample has no nearest neighbors, this is known as the curse of dimensionality.

- The variance of KNN regression is very large

Classification problems

Recall: Supervised learning with a *qualitative* or *categorical* response. As common (if not more) than regression

- *Medical diagnosis*: Given the symptoms a patient shows, predict which of 3 conditions they are attributed to
- *Online banking*: Determine whether a transaction is fraudulent or not, on the basis of the IP address, client's history, etc.
- *Web searching*: Based on a user's attributes and the string of a web search, predict which link a person will click
- *Online advertising*: Predict whether a user will click on an ad

Classification problems

Recall: In classification, the function f_0 we care about is

$$f_0 \triangleq \mathbb{P}_0[Y = y|X_1, X_2, ..., X_p] \tag{12}$$

To get a prediction, we use the Bayes Classifier:

$$\hat{y} = \arg\max_y \mathbb{P}_0[Y = y|X_1, X_2, ..., X_p] \tag{13}$$

Classification problems

Recall: In classification, the function f_0 we care about is

$$f_0 \triangleq \mathbb{P}_0[Y = y|X_1, X_2, ..., X_p] \tag{12}$$

To get a prediction, we use the Bayes Classifier:

$$\hat{y} = \arg\max_y \mathbb{P}_0[Y = y|X_1, X_2, ..., X_p] \tag{13}$$

Example: Suppose $Y \in \{0, 1\}$. We *could* use linear model:

$$\mathbb{P}[Y = 1|\mathbf{X}] = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p \tag{14}$$

Problems:

- This would allow probabilities < 0 and > 1
- Difficult to extend to more than 2 categories

An idea:
Let's apply a function to the result to keep it within $[0, 1]$

$$g^{-1}(z) = \frac{1}{1 + exp(-z)} \tag{15}$$

i.e.

$$\mathbb{P}[Y = 1|\mathbf{X}] = \frac{1}{1 + exp(-(\beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p))} \tag{16}$$

Logistic regression

An idea:
Let's apply a function to the result to keep it within $[0, 1]$

$$g^{-1}(z) = \frac{1}{1 + exp(-z)} \tag{15}$$

i.e.

$$\mathbb{P}[Y = 1|\mathbf{X}] = \frac{1}{1 + exp(-(\beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p))} \tag{16}$$

This is equivalent to modeling the log-odds, e.g.

$$\log \left[\frac{\mathbb{P}[Y = 1|\mathbf{X}]}{\mathbb{P}[Y = 0|\mathbf{X}]} \right] = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p \tag{17}$$

n.b. $exp(\beta_j)$ is commonly referred to as the *odds-ratio* for X_j

Fitting a logistic regression

Let $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), ..., (\mathbf{x}_n, y_n)$ be our training data.

In the linear model

$$\log \left[\frac{\mathbb{P}[Y = 1|\mathbf{X}]}{\mathbb{P}[Y = 0|\mathbf{X}]} \right] = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p, \tag{18}$$

We don't actually observe the left side

- We observe $Y \in \{0, 1\}$, not probabilities
- This prevents us from using e.g. least squares to estimate our parameters

Fitting a logistic regression

Solution:

Let's try to maximize the probability of our training data

Fitting a logistic regression

Solution:
Let's try to maximize the probability of our training data

$$\mathcal{L}(\boldsymbol{\theta}) = \prod_{i=1}^{n} \mathbb{P}(Y = y_i | \mathbf{X} = \mathbf{x}_i) \tag{19}$$

$$= \prod_{i=1}^{n} p_i^{y_i} \cdot (1 - p_i)^{1-y_i} \tag{20}$$

where $p_i = g^{-1}(\beta_0 + \beta_1 x_{1,i} + \cdots + \beta_p x_{p,i})$

Fitting a logistic regression

Solution:
Let's try to maximize the probability of our training data

$$\mathcal{L}(\boldsymbol{\theta}) = \prod_{i=1}^{n} \mathbb{P}(Y = y_i | \mathbf{X} = \mathbf{x}_i) \tag{19}$$

$$= \prod_{i=1}^{n} p_i^{y_i} \cdot (1 - p_i)^{1 - y_i} \tag{20}$$

where $p_i = g^{-1}(\beta_0 + \beta_1 x_{1,i} + \cdots + \beta_p x_{p,i})$

- We look for θ such that $\mathcal{L}(\boldsymbol{\theta})$ is maximized
- aka Maximum likelihood estimation (MLE)
- Has no closed form solution, so solved with numerical methods (e.g. Newton's method)

Fitting a logistic regression

Note:
We typically deal with the log-likelihood:

$$\begin{aligned} \ell(\boldsymbol{\theta}) &= \log \mathcal{L}(\boldsymbol{\theta}) && (21) \\ &= \sum_{i=1}^{n} y_i \log(p_i) + (1 - y_i) \log(1 - p_i) && (22) \\ &= \sum_{k=0}^{1} \sum_{i=1}^{n} \mathbb{I}(Y_i = k) \log(\mathbb{P}(k|\mathbf{X} = \mathbf{x}_i)) && (23) \end{aligned}$$

Fitting a logistic regression

Given our loss function:

$$\ell(\boldsymbol{\theta}) = \sum_{i=1}^{n} y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \tag{24}$$

Fitting a logistic regression

Given our loss function:

$$\ell(\boldsymbol{\theta}) = \sum_{i=1}^{n} y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \tag{24}$$

Where

$$p_i = \frac{1}{1 + \exp(-Z_i)} \tag{25}$$

$$Z_i = \mathbf{X_i}\beta \tag{26}$$

Fitting a logistic regression

Given our loss function:

$$\ell(\boldsymbol{\theta}) = \sum_{i=1}^{n} y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \tag{24}$$

Where

$$p_i = \frac{1}{1 + \exp(-Z_i)} \tag{25}$$

$$Z_i = \mathbf{X_i}\beta \tag{26}$$

We can deriving the gradient using the chain rule:

$$\frac{\partial \ell(\boldsymbol{\theta})}{\partial \beta} = \frac{\partial \ell(\boldsymbol{\theta})}{\partial p_i} \times \frac{\partial p_i}{\partial Z_i} \times \frac{\partial Z_i}{\partial \beta} \tag{27}$$

Logistic regression in R

Estimating uncertainty

- We can estimate the Standard Error of each coefficient (e.g. using Fisher's information)

$$\mathbf{I}_{\mathbf{Y}}(\beta) = -\mathbb{E}_{\beta}[\nabla^2 \ell(\beta)] \tag{28}$$

- The z-statistic (for logistic regression) is the equivalent of the t-statistic (in linear regression):

$$z = \frac{\hat{\beta}_j}{\hat{SE}(\hat{\beta}_j)} \tag{29}$$

- The p-values are test of the null hypothesis $\beta_j = 0$

Logistic regression in R

Example fit

```
> glm.fit=glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume,
    data=Smarket,family=binomial)
> summary(glm.fit)

Call:
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5
    + Volume, family = binomial, data = Smarket)

Deviance Residuals:
   Min      1Q  Median      3Q     Max
 -1.45   -1.20    1.07    1.15    1.33

Coefficients:
             Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.12600    0.24074   -0.52     0.60
Lag1        -0.07307    0.05017   -1.46     0.15
Lag2        -0.04230    0.05009   -0.84     0.40
Lag3         0.01109    0.04994    0.22     0.82
Lag4         0.00936    0.04997    0.19     0.85
Lag5         0.01031    0.04951    0.21     0.83
Volume       0.13544    0.15836    0.86     0.39
```

Example: Predicting credit card default

Predictors:

- `student`: 1 if student, 0 otherwise
- `balance`: credit card balance
- `income`: person's income

In this dataset there is *confounding*, but little collinearity

- Students tend to have higher balances. So, balance is explained by student, but not very well.
- People with a high balance are more likely to default.
- Among people with a given balance, students are less likely to default.

Example: Predicting credit card default

Predictors:

- `student`: 1 if student, 0 otherwise
- `balance`: credit card balance
- `income`: person's income

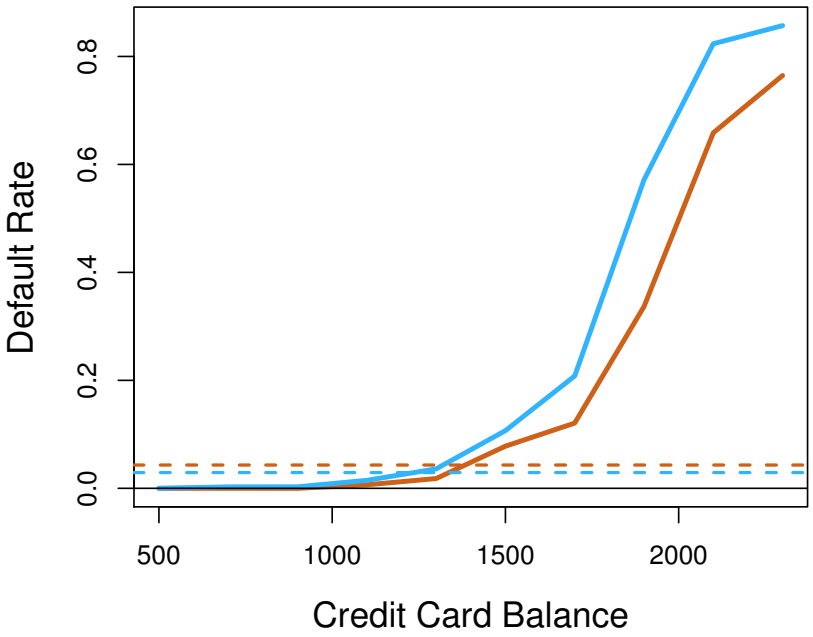

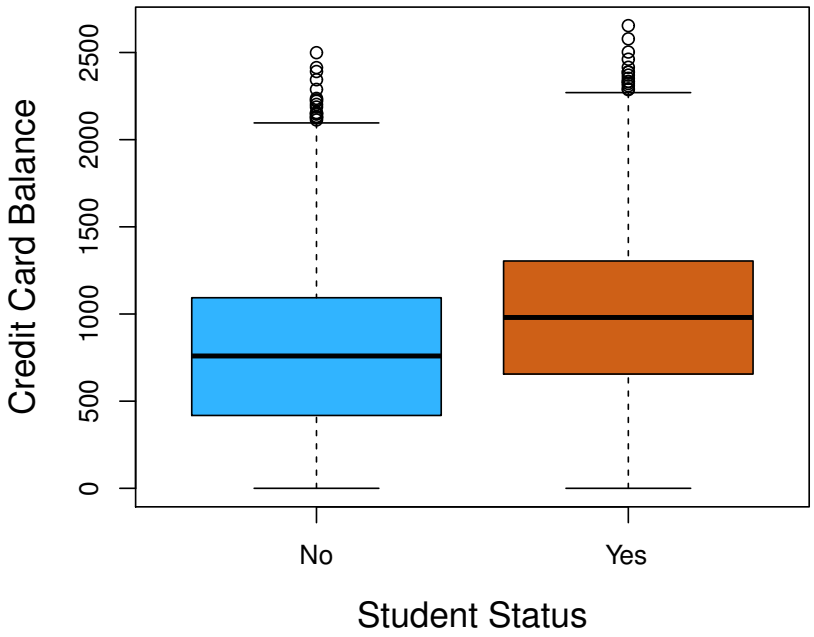

Example: Predicting credit card default

Logistic regression using only `balance`:

	Coefficient	Std. error	Z-statistic	P-value
`Intercept`	−10.6513	0.3612	−29.5	<0.0001
`balance`	0.0055	0.0002	24.9	<0.0001

Logistic regression using only `student`:

	Coefficient	Std. error	Z-statistic	P-value
`Intercept`	−3.5041	0.0707	−49.55	<0.0001
`student[Yes]`	0.4049	0.1150	3.52	0.0004

Logistic regression using all 3 predictors:

	Coefficient	Std. error	Z-statistic	P-value
`Intercept`	−10.8690	0.4923	−22.08	<0.0001
`balance`	0.0057	0.0002	24.74	<0.0001
`income`	0.0030	0.0082	0.37	0.7115
`student[Yes]`	−0.6468	0.2362	−2.74	0.0062

Some issues with logistic regression

- The coefficients become unstable when there is collinearity
  - This also affects the convergence of the fitting algorithm
- When the classes are well separated, the coefficients become unstable
  - This is always the case when $p \geq n - 1$.
- Sometimes may not converge
  - e.g. Needs more iterations

Linear Discriminant Analysis

A linear model (like logistic regression). Unlike logistic regression:

- Does not become unstable when classes are well separated
- With small n and $\mathbf{X}$ approximately normal, is stable
- Popular when we have > 2 classes

Linear Discriminant Analysis

A linear model (like logistic regression). Unlike logistic regression:

- Does not become unstable when classes are well separated
- With small n and $\mathbf{X}$ approximately normal, is stable
- Popular when we have > 2 classes

High level idea:
Model distribution of $\mathbf{X}$ given Y, and apply Bayes' theorem, i.e.

$$\mathbb{P}(Y = k|\mathbf{X} = \mathbf{x}) = \frac{\pi_k f_k(\mathbf{x})}{\sum_{l=1}^{K} \pi_l f_l(\mathbf{x})} \tag{30}$$

- A common assumption is $f_k(\mathbf{x})$ is Gaussian

Linear Discriminant Analysis

Example: $K = 2$ with Gaussian $f_k(\mathbf{x})$ and common σ^2

$$\mathbb{P}(Y = k|\mathbf{X} = \mathbf{x}) = \frac{\pi_k f_k(\mathbf{x})}{\sum_{l=1}^{K} \pi_l f_l(\mathbf{x})} \tag{31}$$

$$= \frac{\pi_k \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(x - \mu_k)^2\right)}{\sum_{l=1}^{2} \pi_l \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(x - \mu_l)^2\right)} \tag{32}$$

Linear Discriminant Analysis

Example: $K = 2$ with Gaussian $f_k(\mathbf{x})$ and common σ^2

$$\mathbb{P}(Y = k|\mathbf{X} = \mathbf{x}) = \frac{\pi_k f_k(\mathbf{x})}{\sum_{l=1}^{K} \pi_l f_l(\mathbf{x})} \tag{31}$$

$$= \frac{\pi_k \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(x - \mu_k)^2\right)}{\sum_{l=1}^{2} \pi_l \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(x - \mu_l)^2\right)} \tag{32}$$

Taking the log and rearranging gives:

$$\delta_k(x) = x \cdot \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k) \tag{33}$$

Linear Discriminant Analysis

Example: $K = 2$ with Gaussian $f_k(\mathbf{x})$ and common σ^2

$$\mathbb{P}(Y = k|\mathbf{X} = \mathbf{x}) = \frac{\pi_k f_k(\mathbf{x})}{\sum_{l=1}^{K} \pi_l f_l(\mathbf{x})} \tag{31}$$

$$= \frac{\pi_k \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(x - \mu_k)^2\right)}{\sum_{l=1}^{2} \pi_l \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(x - \mu_l)^2\right)} \tag{32}$$

Taking the log and rearranging gives:

$$\delta_k(x) = x \cdot \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k) \tag{33}$$

If $\pi_1 = \pi_2$, our Bayes Classifier is:

$$2x(\mu_1 - \mu_2) > \mu_1^2 - \mu_2^2 \tag{34}$$

Example of LDA

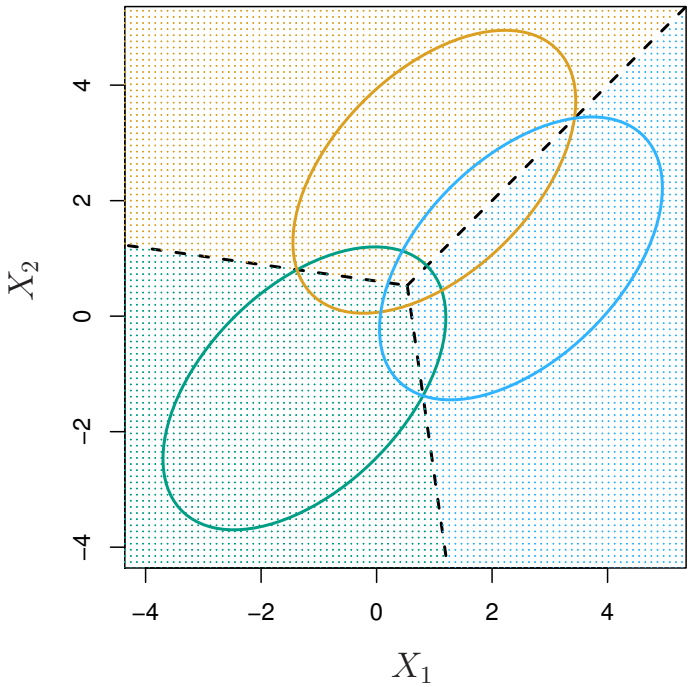

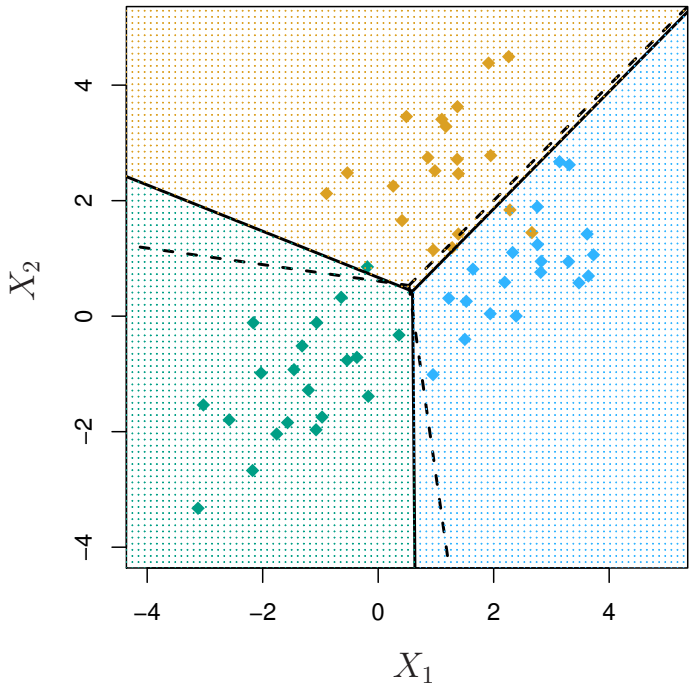

Quadratic Discriminant Analysis

Similar to LDA

- Assumes Gaussian $f_k(\mathbf{x})$
- Unlike LDA:
  - Assumes each class has its own covariance matrix ($\mathbf{\Sigma}_k$)

Quadratic Discriminant Analysis

Similar to LDA

- Assumes Gaussian $f_k(\mathbf{x})$
- Unlike LDA:
  - Assumes each class has its own covariance matrix ($\mathbf{\Sigma}_k$)

This results in a quadratic discriminant function:

$$\begin{aligned}\delta_k(x) &= -\frac{1}{2}(x-\mu_k)^\top \mathbf{\Sigma}_k^{-1}(x-\mu_k) - \frac{1}{2}\log|\mathbf{\Sigma}_k| + \log\pi_k \\ &= -\frac{1}{2}x^\top \mathbf{\Sigma}_k^{-1}x + x^\top \mathbf{\Sigma}_k^{-1}\mu_k - \frac{1}{2}\mu_k^\top \mathbf{\Sigma}_k^{-1}\mu_k - \frac{1}{2}\log|\mathbf{\Sigma}_k| + \log\pi_k \end{aligned} \tag{35}$$

Quadratic Discriminant Analysis

Similar to LDA

- Assumes Gaussian $f_k(\mathbf{x})$
- Unlike LDA:
  - Assumes each class has its own covariance matrix ($\mathbf{\Sigma}_k$)

This results in a quadratic discriminant function:

$$\begin{aligned}\delta_k(x) &= -\frac{1}{2}(x-\mu_k)^\top \mathbf{\Sigma}_k^{-1}(x-\mu_k) - \frac{1}{2}\log|\mathbf{\Sigma}_k| + \log\pi_k \\ &= -\frac{1}{2}x^\top \mathbf{\Sigma}_k^{-1}x + x^\top \mathbf{\Sigma}_k^{-1}\mu_k - \frac{1}{2}\mu_k^\top \mathbf{\Sigma}_k^{-1}\mu_k - \frac{1}{2}\log|\mathbf{\Sigma}_k| + \log\pi_k \end{aligned} \tag{35}$$

This results in more parameters to fit:

- LDA: Kp parameters
- QDA: $Kp(p+1)/2$ parameters

Regularized Discriminant Analysis

Rather than sticking with just LDA or QDA, we can have a combo:

$$\hat{\Sigma}_k(\alpha) = \alpha\hat{\Sigma}_k + (1-\alpha)\hat{\Sigma} \tag{36}$$

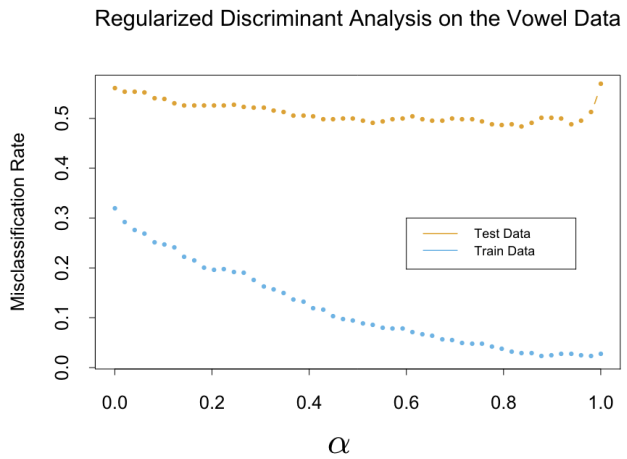

FIGURE 4.7. *Test and training errors for the vowel data, using regularized discriminant analysis with a series of values of* $\alpha \in [0, 1]$. *The optimum for the test data occurs around* $\alpha = 0.9$, *close to quadratic discriminant analysis.*

Linear Discriminant Analysis vs Logistic regression

Assume a two-class setting with one predictor

Linear Discriminant Analysis:

$$\log\left[\frac{p_1(x)}{1-p_1(x)}\right] = c_0 + c_1 x \tag{37}$$

- c_0 and c_1 computed using $\hat{\mu}_0$, $\hat{\mu}_1$, and $\hat{\sigma}2$

Logistic regression:

$$\log\left[\frac{\mathbb{P}[Y=1|x]}{1-\mathbb{P}[Y=1|x]}\right] = \beta_0 + \beta_1 x \tag{38}$$

- β_0 and β_1 estimated using MLE

Comparison of classification methods

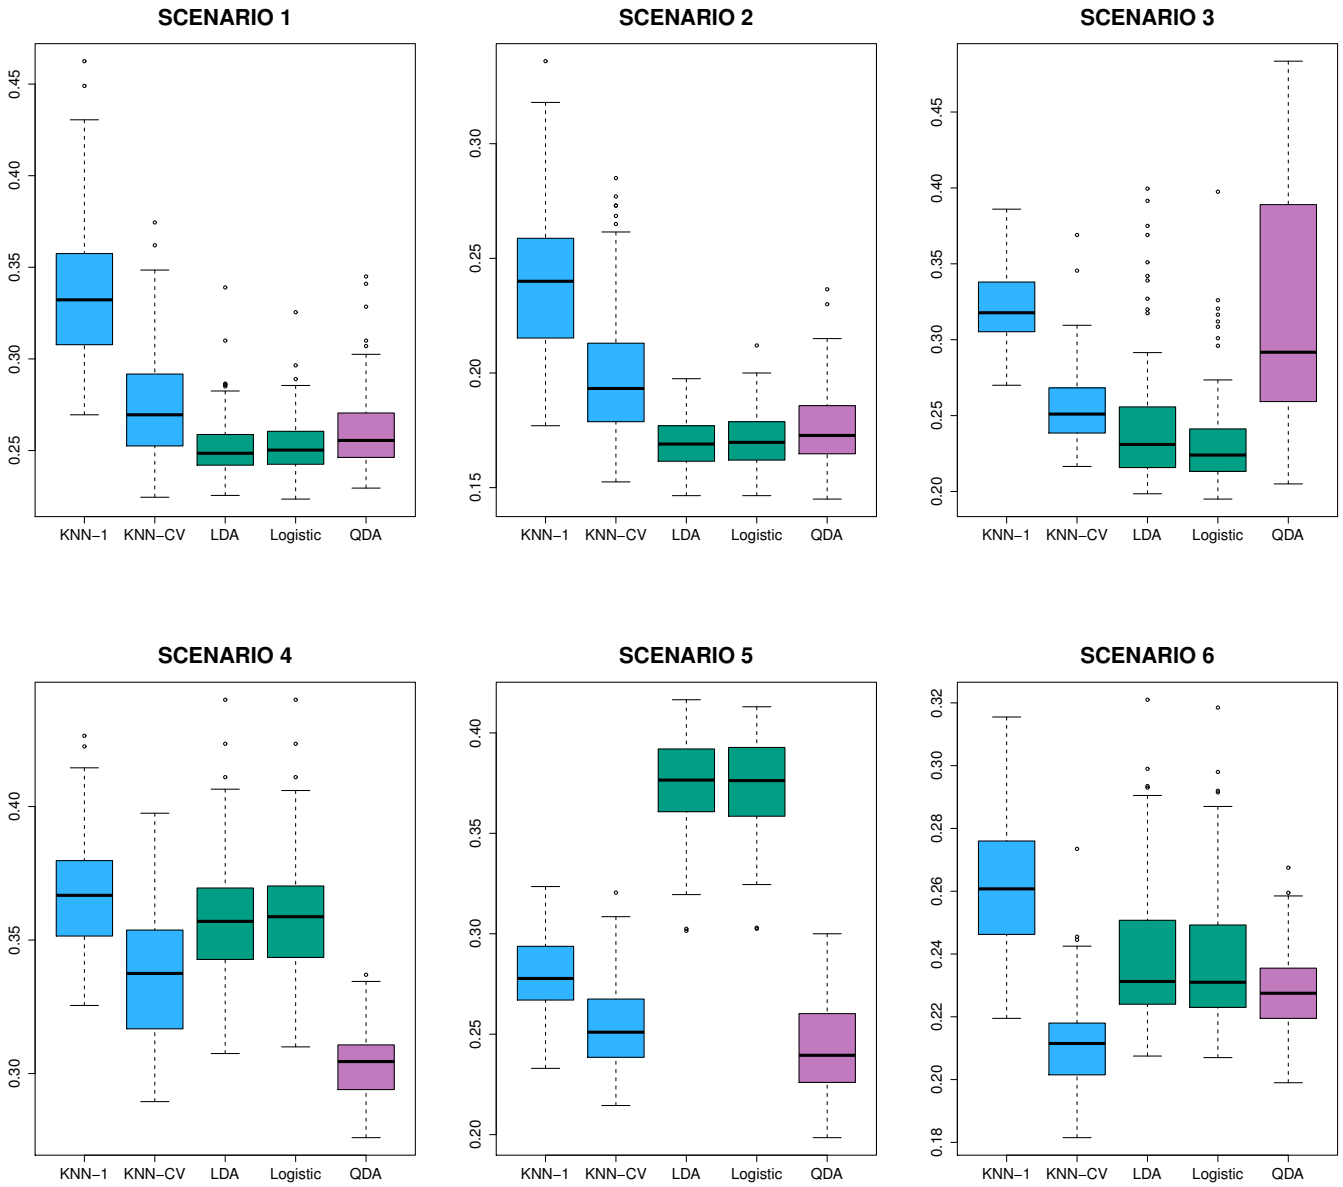

References

[1] ISL. Chapters 3-4.

[2] ESL. Chapters 3.