

Lecture 11: Neural networks

STATS 202: Statistical Learning and Data Science

Linh Tran

stat202@gmail.com



Department of Statistics
Stanford University

July 30, 2025



- ▶ Kaggle predictions due in 11 days
- ▶ Homework 4 is due in 7 days
- ▶ Final exam covers all material up to today
- ▶ No review session this Friday.




- ▶ Introduction
- ▶ Logistic regression
- ▶ Back propagation
- ▶ Function approximation
- ▶ Feature extraction
- ▶ Model generalization
- ▶ Advanced topics



Currently, the most popular algorithm amongst ML practitioners.

- ▶ Many times, used within the context of Artificial Intelligence.
- ▶ Simply a general function estimation algorithm.
- ▶ Though is often hyped up the media.




Gartner's hype cycle for 2018.


Lots of buzz words, but what do they mean?

Definitions:


- ▶ **AI**: human-like machines or programs.
- ▶ **ML**: Algorithms that learn from data.
- ▶ **DL**: A type of ML algorithm, using neural networks (typically with many layers).



But: what exactly are neural networks?




But: what exactly are neural networks?



Some potential answers

- A *universal function approximator*.


But: what exactly are neural networks?



Some potential answers

- ▶ A *universal function approximator*.
- ▶ A *feature extractor*.

But: what exactly are neural networks?



Some potential answers

- ▶ A *universal function approximator*.
- ▶ A *feature extractor*.
- ▶ A *model generalizer*.



Recall: logistic regression is a linear model with a logit link function, i.e.

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p) : \sigma(z) = \frac{1}{1 + \exp(-z)} \quad (1)$$

Let's rephrase this by:

1. Using b to denote β_0 (aka the bias)
2. Using \mathbf{W} to denote $(\beta_0, \dots, \beta_p)$ (aka the weights)
3. Using matrix notation

$$\mathbb{P}(Y = 1|\mathbf{X}) = \underbrace{\sigma}_{\text{non-linearity}}(\mathbf{X}\mathbf{W} + b) \quad (2)$$



When the function is non-linear, our prior option was to do feature transformations, e.g.

- ▶ Expand predictor set (e.g. non-linear transformations, interactions, etc.).
- ▶ Define a kernel (e.g. find a function $f(\cdot, \cdot)$ that is positive definite).



When the function is non-linear, our prior option was to do feature transformations, e.g.

- ▶ Expand predictor set (e.g. non-linear transformations, interactions, etc.).
- ▶ Define a kernel (e.g. find a function $f(\cdot, \cdot)$ that is positive definite).

Another option: build the non-linearity into the model specification, e.g.

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\sigma(\mathbf{X}\mathbf{W}_1 + b_1)\mathbf{W}_2 + b_2) \quad (3)$$



When the function is non-linear, our prior option was to do feature transformations, e.g.

- ▶ Expand predictor set (e.g. non-linear transformations, interactions, etc.).
- ▶ Define a kernel (e.g. find a function $f(\cdot, \cdot)$ that is positive definite).

Another option: build the non-linearity into the model specification, e.g.

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\sigma(\mathbf{X}\mathbf{W}_1 + b_1)\mathbf{W}_2 + b_2) \quad (3)$$

This is a neural network (with 1 hidden layer)!



For logistic regression:

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\underbrace{\mathbf{X} \mathbf{W}}_{p \times 1} + \underbrace{b}_{1 \times 1}) \quad (4)$$



For logistic regression:

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\underbrace{\mathbf{X} \mathbf{W}}_{p \times 1} + \underbrace{b}_{1 \times 1}) \quad (4)$$

For neural networks:

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\underbrace{\sigma(\mathbf{X} \underbrace{\mathbf{W}_1}_{p \times M} + \underbrace{b_1}_{1 \times M})}_{M \times 1} \underbrace{\mathbf{W}_2}_{M \times 1} + \underbrace{b_2}_{1 \times 1}) \quad (5)$$

M specifies how many hidden nodes we have

- ▶ Called '*hidden*' since it's not directly observed by us.
- ▶ Also referred to as '*embeddings*'.



$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\underbrace{\sigma(\mathbf{X} \mathbf{W}_1 + b_1)}_{p \times M} \underbrace{\mathbf{W}_2 + b_2}_{M \times 1}) \quad (6)$$




FIGURE 11.2. Schematic of a single hidden layer, feed-forward neural network.



We can iteratively apply our non-linear operations, e.g.

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\cdots \sigma(\sigma(\mathbf{X}\mathbf{W}_1 + b_1)\mathbf{W}_2 + b_2) \cdots \mathbf{W}_B + b_B) \quad (7)$$

Where B is the number of iterations (i.e. *hidden layers*).



We can iteratively apply our non-linear operations, e.g.

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\cdots \sigma(\sigma(\mathbf{X}\mathbf{W}_1 + b_1)\mathbf{W}_2 + b_2) \cdots \mathbf{W}_B + b_B) \quad (7)$$

Where B is the number of iterations (i.e. *hidden layers*).


n.b. Consequently, we have three hyper-parameters:

- ▶ The number of hidden layers
- ▶ The number of nodes within each layer
- ▶ The activation function





Considerations

1. Increasing both requires more parameters (and can overfit)
2. Deeper networks capture more “complex” functions; wider networks capture more “diverse” features
3. Wider networks result in more “localization”
4. Deeper networks are harder to train (vanishing gradient)
5. In practice, typically start with random sets of hyperparameters and fine tune from there.

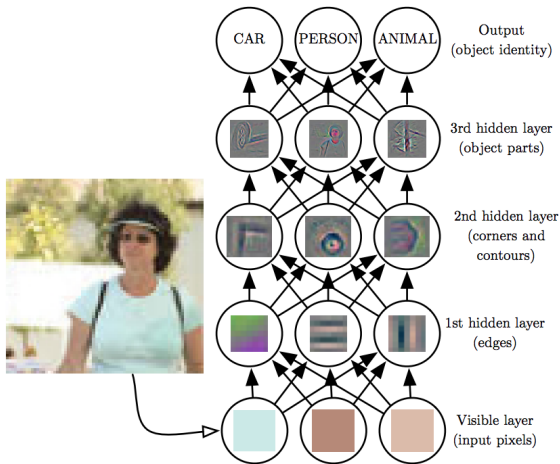


Example 1



Example 2

Deeper networks





Our examples have been for binary outcomes so far.

Question: What about multinomial outcomes

- ▶ e.g. Which of digits 0 through 9 is this photo?



Our examples have been for binary outcomes so far.

Question: What about multinomial outcomes

- ▶ e.g. Which of digits 0 through 9 is this photo?

Recall: for logistic regression, we're modeling

$$\log \left[\frac{\mathbb{P}(Y = 1|\mathbf{X})}{1 - \mathbb{P}(Y = 1|\mathbf{X})} \right] = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p \quad (8)$$

$$= \mathbf{XW} \quad (9)$$

where \mathbf{X} is our $n \times p$ design matrix and \mathbf{W} is our $p \times 1$ parameter vector.



For multinomial regression, let $Y \in \{1, \dots, K\}$. We can model

$$\log \left[\frac{\mathbb{P}(Y = 1|\mathbf{X})}{\mathbb{P}(Y = K|\mathbf{X})} \right] = \mathbf{X}\mathbf{W}_1 \quad (10)$$

$$\log \left[\frac{\mathbb{P}(Y = 2|\mathbf{X})}{\mathbb{P}(Y = K|\mathbf{X})} \right] = \mathbf{X}\mathbf{W}_2 \quad (11)$$

$$\dots = \dots \quad (12)$$

$$\log \left[\frac{\mathbb{P}(Y = K - 1|\mathbf{X})}{\mathbb{P}(Y = K|\mathbf{X})} \right] = \mathbf{X}\mathbf{W}_{K-1} \quad (13)$$

where each \mathbf{W}_k is a $p \times 1$ parameter vector.

Exponentiating both sides and solving for $\mathbb{P}(Y = K|\mathbf{X})$ (using the fact that the probabilities have to sum to 1) gives us

$$\mathbb{P}(Y = K|\mathbf{X}) = \frac{1}{1 + \sum_{k=1}^{K-1} e^{\mathbf{X}\mathbf{W}_k}} \quad (14)$$



Equivalently, can represent the multinomial logistic model as

$$\log \mathbb{P}(Y = 1|\mathbf{X}) = \mathbf{X}\mathbf{W}_1 - \log(Z) \quad (15)$$

$$\log \mathbb{P}(Y = 2|\mathbf{X}) = \mathbf{X}\mathbf{W}_2 - \log(Z) \quad (16)$$

$$\dots = \dots \quad (17)$$

$$\log \mathbb{P}(Y = K|\mathbf{X}) = \mathbf{X}\mathbf{W}_K - \log(Z) \quad (18)$$

resulting in the following probabilities

$$\mathbb{P}(Y = 1|\mathbf{X}) = \frac{\exp(\mathbf{X}\mathbf{W}_1)}{\sum_{k=1}^K \exp(\mathbf{X}\mathbf{W}_k)} \quad (19)$$

$$\mathbb{P}(Y = 2|\mathbf{X}) = \frac{\exp(\mathbf{X}\mathbf{W}_2)}{\sum_{k=1}^K \exp(\mathbf{X}\mathbf{W}_k)} \quad (20)$$

$$\dots = \dots \quad (21)$$

$$\mathbb{P}(Y = K|\mathbf{X}) = \frac{\exp(\mathbf{X}\mathbf{W}_K)}{\sum_{k=1}^K \exp(\mathbf{X}\mathbf{W}_k)} \quad (22)$$



This leads us to the softmax function, i.e.

$$\text{softmax}(\mathbf{XW}_1, \dots, \mathbf{XW}_K)_k = \frac{e^{\mathbf{XW}_k}}{\sum_{l=1}^K e^{\mathbf{XW}_l}} \quad (23)$$

Or, more succinctly, we have

$$\text{softmax}(\mathbf{XW}^K)_k = \frac{\exp((\mathbf{XW}^K)_{k\cdot})}{\sum_{k=1}^K \exp((\mathbf{XW}^K)_{k\cdot})} \quad (24)$$

where the $p \times K$ matrix \mathbf{W}^K is simply the (concatenated) matrix of $\mathbf{W}_1, \dots, \mathbf{W}_K$.

This is what multiclass neural networks are modeling!



Recall: In logistic regression we try to maximize the likelihood

- Equivalent to minimizing the cross-entropy

$$L(y_i, f(\mathbf{X}_i)) = -y_i \log(p_i) - (1 - y_i) \log(1 - p_i), \text{ where} \quad (25)$$

$$p_i = \frac{1}{1 + \exp(-Z_i)} \quad (26)$$

$$Z_i = \mathbf{X}_i \mathbf{W} \quad (27)$$



Recall: In logistic regression we try to maximize the likelihood

- Equivalent to minimizing the cross-entropy

$$L(y_i, f(\mathbf{X}_i)) = -y_i \log(p_i) - (1 - y_i) \log(1 - p_i), \text{ where} \quad (25)$$

$$p_i = \frac{1}{1 + \exp(-Z_i)} \quad (26)$$

$$Z_i = \mathbf{X}_i \mathbf{W} \quad (27)$$

Can apply *the derivative chain rule* to get our gradient, i.e.

$$\frac{\partial L(y_i, f(\mathbf{X}_i))}{\partial \mathbf{W}} = \frac{\partial L(y_i, f(\mathbf{X}_i))}{\partial p_i} \times \frac{\partial p_i}{\partial Z_i} \times \frac{\partial Z_i}{\partial \mathbf{W}} \quad (28)$$



Recall: In logistic regression we try to maximize the likelihood

- Equivalent to minimizing the cross-entropy

$$L(y_i, f(\mathbf{X}_i)) = -y_i \log(p_i) - (1 - y_i) \log(1 - p_i), \text{ where} \quad (25)$$

$$p_i = \frac{1}{1 + \exp(-Z_i)} \quad (26)$$

$$Z_i = \mathbf{X}_i \mathbf{W} \quad (27)$$

Can apply *the derivative chain rule* to get our gradient, i.e.

$$\frac{\partial L(y_i, f(\mathbf{X}_i))}{\partial \mathbf{W}} = \frac{\partial L(y_i, f(\mathbf{X}_i))}{\partial p_i} \times \frac{\partial p_i}{\partial Z_i} \times \frac{\partial Z_i}{\partial \mathbf{W}} \quad (28)$$

Which gives us

$$\frac{\partial L(y_i, f(\mathbf{X}_i))}{\partial \mathbf{W}} = \mathbf{X}_i (y_i - p_i) \quad (29)$$



Neural networks are simply a generalization of the logistic regression case, e.g. for

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\sigma(\mathbf{X} \mathbf{W}_1) \mathbf{W}_2) \quad (30)$$



Neural networks are simply a generalization of the logistic regression case, e.g. for

$$\mathbb{P}(Y = 1|\mathbf{X}) = \sigma(\sigma(\mathbf{X}\mathbf{W}_1)\mathbf{W}_2) \quad (30)$$

Our loss is

$$L(y_i, f(\mathbf{X}_i)) = -y_i \log(p_i) - (1 - y_i) \log(1 - p_i), \text{ where} \quad (31)$$

$$p_i = \frac{1}{1 + \exp(-Z_{2,i})} \quad (32)$$

$$Z_{2,i} = h_i \mathbf{W}_2 \quad (33)$$

$$h_i = \frac{1}{1 + \exp(-Z_{1,i})} \quad (34)$$

$$Z_{1,i} = \mathbf{X} \mathbf{W}_1 \quad (35)$$



Our loss is

$$L(y_i, f(\mathbf{X}_i)) = -y_i \log(p_i) - (1 - y_i) \log(1 - p_i), \text{ where} \quad (36)$$

$$p_i = \frac{1}{1 + \exp(-Z_{2,i})} \quad (37)$$

$$Z_{2,i} = h_i \mathbf{W}_2 \quad (38)$$

$$h_i = \frac{1}{1 + \exp(-Z_{1,i})} \quad (39)$$

$$Z_{1,i} = \mathbf{X} \mathbf{W}_1 \quad (40)$$

Applying *the derivative chain rule*:

$$\frac{\partial L(y_i, f(\mathbf{X}_i))}{\partial \mathbf{W}_2} = \frac{\partial L(y_i, f(\mathbf{X}_i))}{\partial p_i} \times \frac{\partial p_i}{\partial Z_{2,i}} \times \frac{\partial Z_{2,i}}{\partial \mathbf{W}_2}$$

$$\frac{\partial L(y_i, f(\mathbf{X}))}{\partial \mathbf{W}_1} = \frac{\partial L(y_i, f(\mathbf{X}))}{\partial p_i} \times \frac{\partial p_i}{\partial Z_{2,i}} \times \frac{\partial Z_{2,i}}{\partial h_i} \times \frac{\partial h_i}{\partial Z_{1,i}} \times \frac{\partial Z_{1,i}}{\partial \mathbf{W}_1}$$



Our gradient is estimated using our data, i.e.
 $(y_i, \mathbf{X}_i) : i = 1, 2, \dots, n$.



Our gradient is estimated using our data, i.e.

$(y_i, \mathbf{X}_i) : i = 1, 2, \dots, n$.

We can estimate it using, e.g.

- ▶ **Stochastic gradient descent**: estimating our (full) gradient using just one observation.
- ▶ **Gradient descent**: estimating our (full) gradient using all observations.
- ▶ **Mini-batch gradient descent**: using a (random) subsample of our observations.

Each will trade off between variance for the gradient and memory size.



Our gradient is estimated using our data, i.e.

$(y_i, \mathbf{X}_i) : i = 1, 2, \dots, n.$

We can estimate it using, e.g.

- ▶ **Stochastic gradient descent:** estimating our (full) gradient using just one observation.
- ▶ **Gradient descent:** estimating our (full) gradient using all observations.
- ▶ **Mini-batch gradient descent:** using a (random) subsample of our observations.

Each will trade off between variance for the gradient and memory size.

When done iteratively, we'll typically specify a stopping point (e.g. by using a dev set).



The use of non-linearities results in multiple minima, tendency to overfit, and can be unstable.

Some considerations to make:

- ▶ Set initial weight values near zero.
- ▶ Over parameterize and regularize heavily.
- ▶ Standardize input features.
- ▶ Use a dev set and stop training earlier.
- ▶ Try out different weight randomizations and take the one with the lowest (validated) error.
 - ▶ Or average the predictions (or apply bagging).



Estimating neural network parameters simply requires '*propagating back*' errors.

- ▶ We're just applying (matrix) multiplications
 - ▶ GPU's can be very good for this
- ▶ Matrix multiplications can get pretty big (for large networks)
 - ▶ Commonly not worth it to use the Hessian
- ▶ Should be careful with large values going into sigmoid activations
 - ▶ Results in saturated gradients



Hornik's theorem

Whenever the activation function is continuous, bounded, and non-constant, then, for arbitrary compact subsets $X \subseteq \mathbb{R}^k$, standard multilayer feedforward networks can approximate any continuous function on X arbitrarily well with respect to uniform distance, provided that sufficiently many hidden units are available.



Hornik's theorem

Whenever the activation function is continuous, bounded, and non-constant, then, for arbitrary compact subsets $X \subseteq \mathbb{R}^k$, standard multilayer feedforward networks can approximate any continuous function on X arbitrarily well with respect to uniform distance, provided that sufficiently many hidden units are available.

In words: A 2-layer neural network with enough hidden nodes can closely approximate any continuous function $f(x)$.

References:


Cybenko (1989) "Approximations by superpositions of sigmoidal function"

Hornik (1991) "Approximation Capabilities of Multilayer Feedforward Networks"

Leshno and Schocken (1993) "Multilayer Feedforward Networks with Non-Polynomial Activation Functions Can Approximate Any Function"



Given enough hidden nodes, we can approximate any function.



Check out this *visual example* of this.



Some caveats to the theorem

- ▶ We're approximating the function within some bound, i.e.
 $|\hat{f}_n(x) - f(x)| < \epsilon.$
- ▶ Result is meant for *continuous* functions on *compact* subsets of \mathbb{R} .
- ▶ Nothing is guaranteed on the how quickly we can learn the function's parameters.
- ▶ Other function estimators also do a good job approximating!



For non-linear functions,

Logistic regression: expand our feature set via transformations

Neural network: define the model non-linearly



For non-linear functions,

Logistic regression: expand our feature set via transformations

- ▶ We have to specify the feature transformations

Neural network: define the model non-linearly

- ▶ The model learns the feature transformations



For non-linear functions,

Logistic regression: expand our feature set via transformations


- ▶ We have to specify the feature transformations

Neural network: define the model non-linearly


- ▶ The model learns the feature transformations
- ▶ *This helps us greatly when dealing with abstract or high dimensional problems (e.g. images & text)!*



How do the feature transformations get learned?



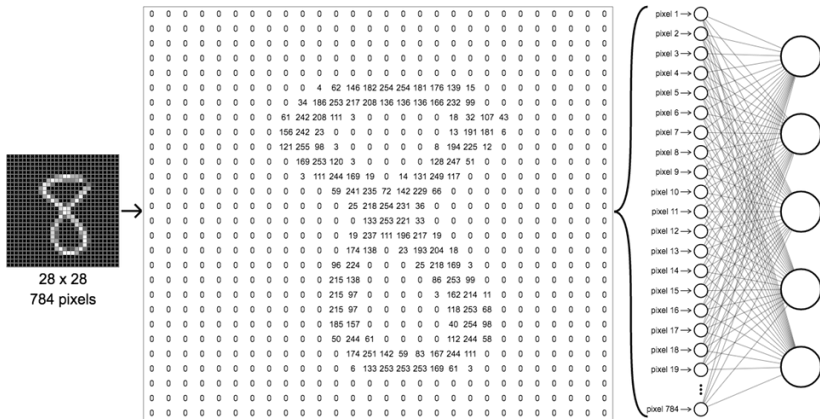
Original representation of curves



Hidden layer representation of curves



Well demonstrated by *Chris Olah's blog*.

Example: The MNIST data



A linear model (e.g. for multinomial logistic regression)


Example: The MNIST data



Neural network's learned (kernel) features.




Knowing that we can simply back propagate errors via multiplication opens many doors for us, e.g.



Google's InceptionNet architecture



Knowing that we can simply back propagate errors via multiplication opens many doors for us, e.g.



Google's InceptionNet architecture

Problem: *large networks are vulnerable to vanishing/exploding gradients.*



Recall: Our gradient is simply a product of partial derivatives.

$$\frac{\partial C}{\partial b_1} = \sigma'(z_1) \times w_2 \times \sigma'(z_2) \times w_3 \times \sigma'(z_3) \times w_4 \times \sigma'(z_4) \times \frac{\partial C}{\partial a_4}$$



Neural Network

Example neural network and gradient.



Recall: Our gradient is simply a product of partial derivatives.

$$\frac{\partial C}{\partial b_1} = \sigma'(z_1) \times w_2 \times \sigma'(z_2) \times w_3 \times \sigma'(z_3) \times w_4 \times \sigma'(z_4) \times \frac{\partial C}{\partial a_4}$$



Neural Network


Example neural network and gradient.

Question: What is the derivative of the sigmoid function?

The vanishing gradient problem




Google's InceptionNet address this via strategically placed loss functions.



Google's InceptionNet architecture



Embeddings: The skip gram model




$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

$$p(w_O | w_I) = \frac{\exp(v'_{w_O} \top v_{w_I})}{\sum_{w=1}^W \exp(v'_w \top v_{w_I})}$$




We can also train models end to end, e.g.



Encoder-decoder architecture




We can also train models end to end, e.g.



Encoder-decoder architecture

Or: in a modular fashion, e.g. pre-training.

Or over multiple tasks (i.e. multi-task learning), e.g.



Kendall et al. 2017's multi-task model

Many researchers will create unique architectures for specific problems, e.g. *Instacart*




The prediction problem

Many researchers will create unique architectures for specific problems, e.g. *Instacart*



The prediction problem



The initial solution

Another example using the Netflix data.


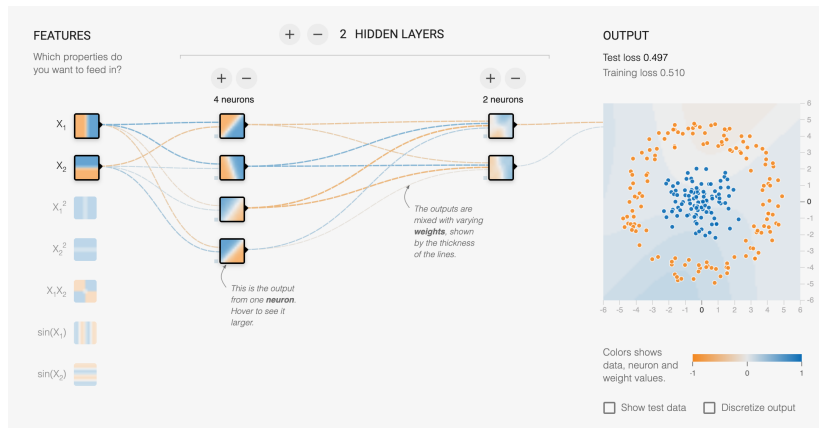


Figure 5. A sample DNN architecture for learning movie embeddings from collaborative filtering data.

The tensorflow playground



An *interactive demo* that allows you to play with a neural network.





Some proposed settings

Core Training Hyperparameters

Hyperparameter	Recommended Range	Notes
Learning Rate	1e-5 to 5e-4	Start with 2e-5 for BERT-like models, 3e-4 for GPT-style
Batch Size	8-64	Use gradient accumulation if memory limited
Number of Epochs	2-10	Early stopping recommended, typically 3-5 epochs
Weight Decay	0.01-0.1	0.01 is common default, increase for overfitting

Learning Rate Scheduling

Hyperparameter	Recommended Range	Notes
Warmup Steps	500-2000	Or 6-10% of total training steps
Warmup Ratio	0.06-0.1	Alternative to fixed warmup steps
LR Scheduler	Linear, Cosine	Linear decay most common for fine-tuning
Min LR Factor	0.0-0.1	For cosine annealing, minimum LR as fraction of max



Some proposed settings

Regularization

Hyperparameter	Recommended Range	Notes
Dropout Rate	0.1-0.3	Often keep pre-trained model's dropout
Attention Dropout	0.1-0.2	Dropout applied to attention weights
Hidden Dropout	0.1-0.3	Dropout in feed-forward layers
Label Smoothing	0.0-0.1	For classification tasks, 0.1 is common

Optimization Settings

Hyperparameter	Recommended Range	Notes
Optimizer	AdamW	Standard choice for transformers
Beta1	0.9	Adam momentum parameter
Beta2	0.98-0.999	0.999 default, 0.98 for some large models
Epsilon	1e-6 to 1e-8	Adam numerical stability parameter
Gradient Clipping	0.5-2.0	Max gradient norm, 1.0 is common



Some proposed settings

Model Size	Learning Rate	Batch Size	Special Notes
Small (< 100M params)	3e-4 to 1e-3	16-64	Can handle higher learning rates
Base (100M-1B params)	1e-5 to 5e-4	8-32	BERT-base, GPT-2 medium
Large (1B+ params)	5e-6 to 1e-4	4-16	May need gradient accumulation

Data Size Considerations

Dataset Size	Epochs	Learning Rate	Notes
Small (< 10K samples)	5-10	Lower end of range	High risk of overfitting
Medium (10K-100K)	3-5	Standard range	Most common scenario
Large (100K+ samples)	2-3	Can use higher LR	Less overfitting risk



Neural net related core topics:

- ▶ Weight initializations
- ▶ Activation functions
- ▶ Optimization functions
- ▶ Loss functions
- ▶ Normalization
- ▶ Regularization / dropout
- ▶ Model architectures
- ▶ Hyperparameter optimization
- ▶ Bayesian neural networks
- ▶ Computation graphs
- ▶ Software / platforms
- ▶ Encoding / adding outside knowledge
- ▶ Hardware accelerators



Neural net applied topics:

- ▶ Computer vision
- ▶ Natural language processing
- ▶ Signal processing
- ▶ Generative models
- ▶ Unsupervised learning
- ▶ Reinforcement learning
- ▶ One/Zero shot learning
- ▶ Transfer learning
- ▶ Auto-ML
- ▶ Memory Augmented Neural Networks



- [1] ESL. Chapter 11
- [2] Pancha N, Zhai A, Leskovec J, Rosenberg C. PinnerFormer: Sequence Modeling for User Representation at Pinterest. arXiv 2022.